

8

Color Images

Color images are involved in every aspect of our lives, where they play an important role in everyday activities such as television, photography, and printing. Color perception is a fascinating and complicated phenomenon that has occupied the interest of scientists, psychologists, philosophers, and artists for hundreds of years [38, 39]. In this chapter, we focus on those technical aspects of color that are most important for working with digital color images. Our emphasis will be on understanding the various representations of color and correctly utilizing them when programming. Additional color-related issues, such as color quantization and colorimetric color spaces, are covered in Volume 2 [6].

8.1 RGB Color Images

The RGB color schema encodes colors as combinations of the three primary colors: red (R), green (G), and blue (B). This scheme is widely used for transmission, representation, and storage of color images on both analog devices such as television sets and digital devices such as computers, digital cameras, and scanners. For this reason, many image-processing and graphics programs use the RGB schema as their internal representation for color images, and most language libraries, including Java's imaging APIs, use it as their standard image representation.

RGB is an *additive* color system, which means that all colors start with black and are created by adding the primary colors. You can think of color formation in this system as occurring in a dark room where you can overlay

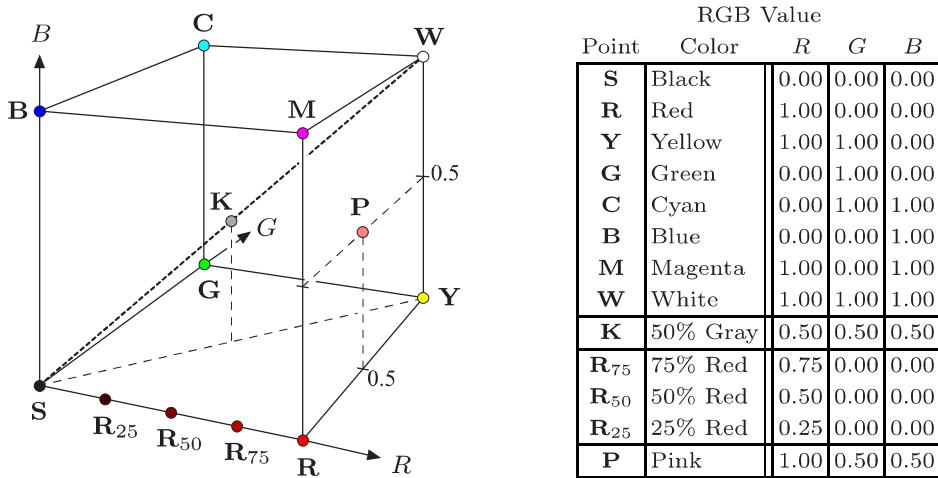


Figure 8.1 Representation of the RGB color space as a three-dimensional unit cube. The primary colors red (R), green (G), and blue (B) form the coordinate system. The “pure” red color (**R**), green (**G**), blue (**B**), cyan (**C**), magenta (**M**), and yellow (**Y**) lie on the vertices of the color cube. All the shades of gray, of which **K** is an example, lie on the diagonal between black **S** and white **W**.

three beams of light—one red, one green, and one blue—on a sheet of white paper. To create different colors, you would modify the intensity of each of these beams independently. The distinct intensity of each primary color beam controls the shade and brightness of the resulting color. The colors gray and white are created by mixing the three primary color beams at the same intensity. A similar operation occurs on the screen of a color television or CRT¹-based computer monitor, where tiny, close-lying dots of red, green, and blue phosphorous are simultaneously excited by a stream of electrons to distinct energy levels (intensities), creating a seemingly continuous color image.

The RGB color space can be visualized as a three-dimensional unit cube in which the three primary colors form the coordinate axis. The RGB values are positive and lie in the range $[0, C_{\max}]$; for most digital images, $C_{\max} = 255$. Every possible color **C**_i corresponds to a point within the RGB color cube of the form

$$\mathbf{C}_i = (R_i, G_i, B_i),$$

where $0 \leq R_i, G_i, B_i \leq C_{\max}$. RGB values are often normalized to the interval $[0, 1]$ so that the resulting color space forms a unit cube (Fig. 8.1). The point **S** = (0, 0, 0) corresponds to the color black, **W** = (1, 1, 1) corresponds to the color white, and all the points lying on the diagonal between **S** and **W** are shades of gray created from equal color components $R = G = B$.

Figure 8.2 shows a color test image and its corresponding RGB color components, displayed here as intensity images. We will refer to this image in a

¹ Cathode ray tube.

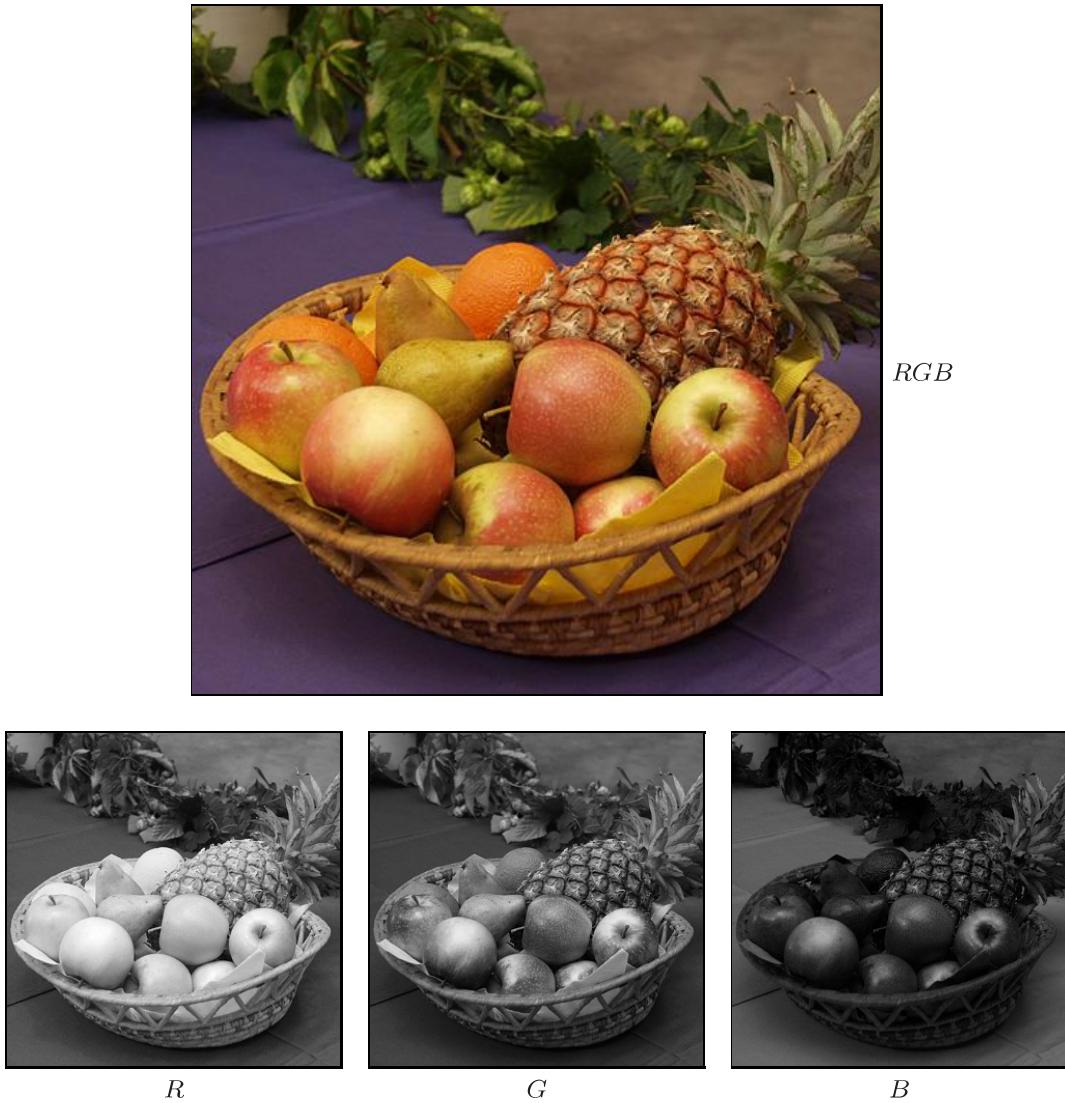


Figure 8.2 A color image and its corresponding RGB channels. The fruits depicted are mainly yellow and red and therefore have high values in the *R* and *G* channels. In these regions, the *B* content is correspondingly lower (represented here by darker gray values) except for the bright highlights on the apple, where the color changes gradually to white. The tabletop in the foreground is purple and therefore displays correspondingly higher values in its *B* channel.

number of examples that follow in this chapter.

RGB is a very simple color system, and as demonstrated in Sec. 8.2, a basic knowledge of it is often sufficient for processing color images or transforming them into other color spaces. At this point, we will not be able to determine what color a particular RGB pixel corresponds to in the real world, or even what the primary colors red, green, and blue truly mean in a physical (i. e., colorimetric) sense. Instead, in this volume we will rely on our intuitive understanding of color and address colorimetry and color spaces in detail in

Vol. 2 [6, Sec. 6].

8.1.1 Organization of Color Images

Color images are represented in the same way as grayscale images, by using an array of pixels in which different models are used to order the individual color components. In the next sections we will examine the difference between *true color* images, which utilize colors uniformly selected from the entire color space, and so-called *palleted* or *indexed* images, in which only a select set of distinct colors are used. Deciding which type of image to use depends on the requirements of the application.

True color images

A pixel in a true color image can represent any color in its color space, as long as it falls within the (discrete) range of its individual color components. True color images are appropriate when the image contains many colors with subtle differences, as occurs in digital photography and photo-realistic computer graphics. Next we look at two methods of ordering the color components in true color images: *component ordering* and *packed ordering*.

Component ordering. In *component ordering* (also referred to as *planar ordering*) the color components are laid out in separate arrays of identical dimensions. In this case, the color image

$$I = (I_R, I_G, I_B)$$

can be thought of as a vector of related intensity images I_R , I_G , and I_B (Fig. 8.3), and the RGB component values of the color image I at position (u, v) are obtained by accessing all three intensity images as follows:

$$\begin{pmatrix} R_{u,v} \\ G_{u,v} \\ B_{u,v} \end{pmatrix} \leftarrow \begin{pmatrix} I_R(u, v) \\ I_G(u, v) \\ I_B(u, v) \end{pmatrix}. \quad (8.1)$$

Packed ordering. In *packed ordering*, the component values that represent the color of a particular pixel are packed together into a single element of the image array (Fig. 8.4) so that

$$I(u, v) = (R_{u,v}, G_{u,v}, B_{u,v}).$$

The RGB value of a packed image I at the location (u, v) is obtained by accessing the individual components of the color pixel as

$$\begin{pmatrix} R_{u,v} \\ G_{u,v} \\ B_{u,v} \end{pmatrix} \leftarrow \begin{pmatrix} \text{Red}(I(u, v)) \\ \text{Green}(I(u, v)) \\ \text{Blue}(I(u, v)) \end{pmatrix}. \quad (8.2)$$

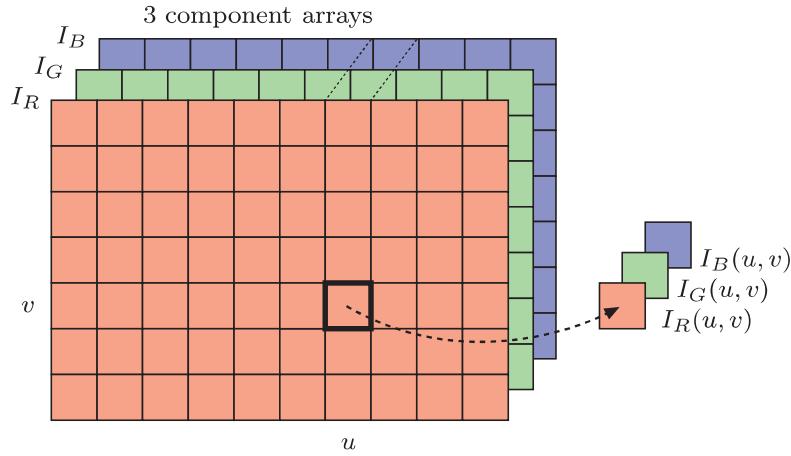


Figure 8.3 RGB color image in component ordering. The three color components are laid out in separate arrays I_R , I_G , I_B of the same size.

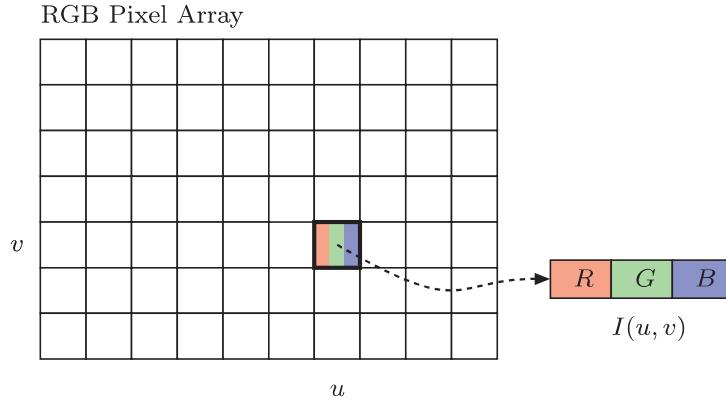


Figure 8.4 RGB-color image using packed ordering. The three color components R , G , and B are placed together in a single array element.

The access functions, `Red()`, `Green()`, `Blue()`, will depend on the specific implementation used for encoding the color pixels.

Indexed images

Indexed images permit only a limited number of distinct colors and therefore are used mostly for illustrations and graphics that contain large regions of the same color. Often these types of images are stored in indexed GIF or PNG files for use on the Web. In these indexed images, the pixel array does not contain color or brightness data but instead consists of integer numbers k that are used to index into a color table or “palette”

$$P(k) = (r_k, g_k, b_k),$$

for $k = 0 \dots N-1$ (Fig. 8.5). N is the size of the color table and therefore also

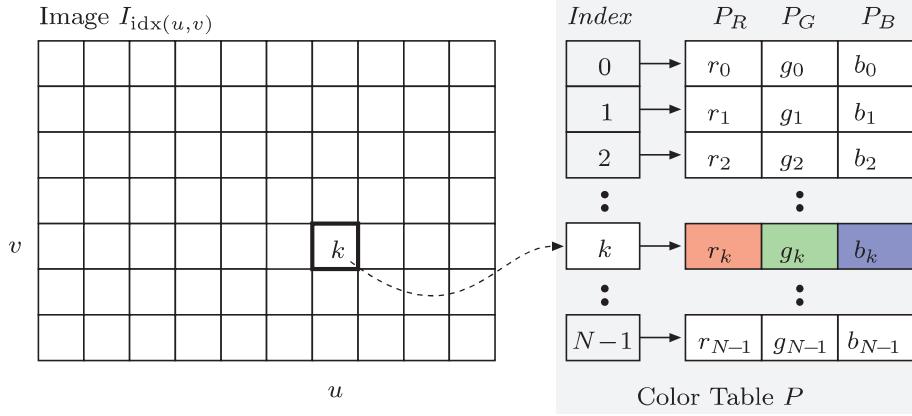


Figure 8.5 RGB indexed image. Instead of a full color value, each pixel in $I_{\text{idx}}(u, v)$ contains an index k . The color value for each k is defined by an entry in the color table or “palette” $P[k]$.

the maximum number of distinct image colors (typically $N = 2$ to 256). Since the color table can contain any RGB color value (r_k, g_k, b_k) , it must be saved as part of the image. The RGB component values of an indexed image I_{idx} at location (u, v) are obtained as

$$\begin{pmatrix} R \\ G \\ B \end{pmatrix} \leftarrow \begin{pmatrix} P_R(k) \\ P_G(k) \\ P_B(k) \end{pmatrix} = \begin{pmatrix} r_k \\ g_k \\ b_k \end{pmatrix}, \quad \text{with } k = I_{\text{idx}}(u, v). \quad (8.3)$$

During the transformation from a true color image to an indexed image (for example, from a JPEG image to a GIF image), the problem of optimal color reduction, or color quantization, arises. Color quantization is the process of determining an optimal color table and then mapping it to the original colors. This process is described in detail in Vol. 2 [6, Sec. 5].

8.1.2 Color Images in ImageJ

ImageJ provides two simple types of color images:

- RGB full-color images (24-bit “RGB color”)
- Indexed images (“8-bit color”)

RGB true color images

RGB color images in ImageJ use a packed order (see Sec. 8.1.1), where each color pixel is represented by a 32-bit `int` value. As Fig. 8.6 illustrates, 8 bits are used to represent each of the RGB components, which limits the range of

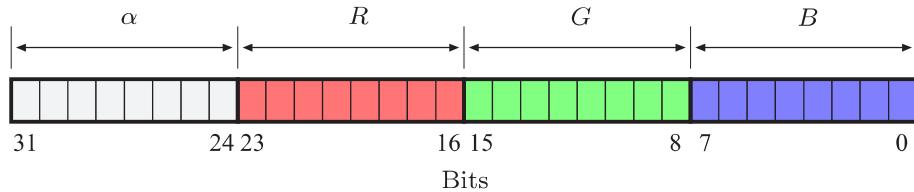


Figure 8.6 Structure of an RGB color pixel in Java. Within a 32-bit `int`, 8 bits are allocated, in the following order, for each of the color components R , G , B as well as the transparency α (unused in ImageJ).

the individual components to 0 to 255. The remaining 8 bits are reserved for the transparency,² or *alpha* (α), component. This is also the usual ordering in Java³ for RGB color images.

Accessing RGB pixel values. RGB color images are represented by an array of pixels, the elements of which are standard Java `ints`. To disassemble the packed `int` value into the three color components, you apply the appropriate bitwise shifting and masking operations. In the following example, we assume that the image processor `ip` contains an RGB color image:

```

1 int c = ip.getPixel(u,v);      // a color pixel
2 int r = (c & 0xff0000) >> 16; // red value
3 int g = (c & 0x00ff00) >> 8; // green value
4 int b = (c & 0x0000ff);     // blue value

```

In this example, each of the RGB components of the packed pixel `c` are isolated using a bitwise AND operation (`&`) with an appropriate bit mask (following convention, bit masks are given in hexadecimal⁴ notation), and afterwards the extracted bits are shifted right by 16 (for R) or 8 (for G) bit positions (see Fig. 8.7).

The “construction” of an RGB pixel from the individual R , G , and B values is done in the opposite direction using the bitwise OR operator (`|`) and shifting the bits left (`<<`):

```

1 int r = 169; // red value
2 int g = 212; // green value
3 int b = 17;  // blue value
4 int c = ((r & 0xff)<<16) | ((g & 0xff)<<8) | b & 0xff;
5 ip.putPixel(u,v,C);

```

Masking the component values with `0xff` works in this case because except for the bits in positions 0 to 7 (values in the range 0 to 255), all the other bits are

² The transparency value α (alpha) represents the ability to see through a color pixel onto the background. At this time, the α channel is unused in ImageJ.

³ Java Advanced Window Toolkit (AWT).

⁴ The mask `0xff0000` is of type `int` and represents the 32-bit binary pattern `00000000111111100000000000000000`.

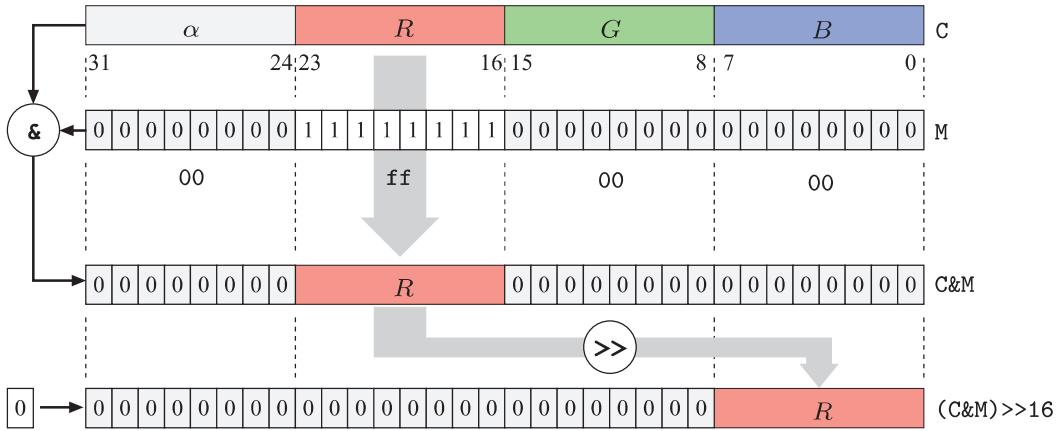


Figure 8.7 Decomposition of a 32-bit RGB color pixel using bit operations. The **R** component (bits 16–23) of the RGB pixels **C** (above) is isolated using a bitwise AND operation ($\&$) together with a bit mask **M** = `0xff0000`. All bits except the **R** component are set to the value 0, while the bit pattern within the **R** component remains unchanged. This bit pattern is subsequently shifted 16 positions to the right ($>>$), so that the **R** component is moved into the lowest 8 bits and its value lies in the range of 0 to 255. During the shift operation, zeros are filled in from the left.

already set to zero. A complete example of manipulating an RGB color image using bit operations is presented in Prog. 8.1. Instead of accessing color pixels using ImageJ's access functions, these programs directly access the pixel array for increased efficiency (see also Sec. B.1.3).

The ImageJ class `ColorProcessor` provides an easy to use alternative which returns the separated RGB components (as an `int` array with three elements). In the following example that demonstrates its use, `ip` is of type `ColorProcessor`:

```

1 int[] RGB = new int[3];
2 ...
3 RGB = ip.getPixel(u,v,RGB);
4 int r = RGB[0];
5 int g = RGB[1];
6 int b = RGB[2];
7 ...
8 ip.putPixel(u,v,RGB);

```

A more detailed and complete example is shown by the simple plugin in Prog. 8.2, which increases the value of all three color components of an RGB image by 10 units. Notice that the plugin limits the resulting component values to 255, because the `putPixel()` method only uses the lowest 8 bits of each component and does not test if the value passed in is out of the permitted 0 to 255 range. Without this test, arithmetic overflow errors can occur. The price for using this access method, instead of direct array access, is a noticeably longer running time (approximately a factor of 4 when compared to the version

```

1 // File Brighten_Rgb_1.java
2
3 import ij.ImagePlus;
4 import ij.plugin.filter.PlugInFilter;
5 import ij.process.ImageProcessor;
6
7 public class Brighten_Rgb_1 implements PlugInFilter {
8
9     public int setup(String arg, ImagePlus im) {
10         return DOES_RGB; // this plugin works on RGB images
11     }
12
13    public void run(ImageProcessor ip) {
14        int[] pixels = (int[]) ip.getPixels();
15
16        for (int i = 0; i < pixels.length; i++) {
17            int c = pixels[i];
18            // split the color pixel into RGB components
19            int r = (c & 0xff0000) >> 16;
20            int g = (c & 0x00ff00) >> 8;
21            int b = (c & 0x0000ff);
22            // modify colors
23            r = r + 10; if (r > 255) r = 255;
24            g = g + 10; if (g > 255) g = 255;
25            b = b + 10; if (b > 255) b = 255;
26            // reassemble the color pixel and insert into pixel array
27            pixels[i]
28                = ((r & 0xff)<<16) | ((g & 0xff)<<8) | b & 0xff;
29        }
30    }
31
32 } // end of class Brighten_Rgb_1

```

Program 8.1 Working with RGB color images using bit operations (ImageJ plugin, version 1). This plugin increases the values of all three color components by 10 units. It demonstrates the use of direct access to the pixel array (line 17), the separation of color components using bit operations (lines 19–21), and the reassembly of color pixels after modification (line 28). The value `DOES_RGB` (defined in the interface `PlugInFilter`) returned by the `setup()` method indicates that this plugin is designed to work on RGB formatted true color images (line 10).

in Prog. 8.1).

Opening and saving RGB images. ImageJ supports the following types of image formats for RGB true color images:

- **TIFF** (only uncompressed): 3×8 -bit RGB. TIFF color images with 16-bit depth are opened as an image stack consisting of three 16-bit intensity images.
- **BMP, JPEG**: 3×8 -bit RGB.
- **PNG**: 3×8 -bit RGB.

```

1 // File Brighten_Rgb_2.java
2
3 import ij.ImagePlus;
4 import ij.plugin.filter.PlugInFilter;
5 import ij.process.ColorProcessor;
6 import ij.process.ImageProcessor;
7
8 public class Brighten_Rgb_2 implements PlugInFilter {
9     static final int R = 0, G = 1, B = 2; // component indices
10
11    public int setup(String arg, ImagePlus im) {
12        return DOES_RGB; // this plugin works on RGB images
13    }
14
15    public void run(ImageProcessor ip) {
16        // make sure the image is of type ColorProcessor
17        ColorProcessor cp = (ColorProcessor) ip;
18        int[] RGB = new int[3];
19
20        for (int v = 0; v < cp.getHeight(); v++) {
21            for (int u = 0; u < cp.getWidth(); u++) {
22                cp.getPixel(u, v, RGB);
23                RGB[R] = Math.min(RGB[R]+10, 255); // add 10 and
24                RGB[G] = Math.min(RGB[G]+10, 255); // limit to 255
25                RGB[B] = Math.min(RGB[B]+10, 255);
26                cp.putPixel(u, v, RGB);
27            }
28        }
29    }
30
31 } // end of class Brighten_Rgb_2

```

Program 8.2 Working with RGB color images without bit operations (ImageJ plugin, version 2). This plugin increases the values of all three color components by 10 units using the access methods `getPixel(int, int, int[])` and `putPixel(int, int, int[])` from the class `ColorProcessor` (lines 22 and 26, respectively). The running time, because of the method calls, is approximately four times higher than that of version 1 (Prog. 8.1).

- **RAW:** using the ImageJ menu `File→Import→Raw`, RGB images can be opened whose format is not directly supported by ImageJ. It is then possible to select different arrangements of the color components.

Creating RGB images. The simplest way to create a new RGB image using ImageJ is to use an instance of the class `ColorProcessor`, as the following example demonstrates:

```

1 int w = 640, h = 480;
2 ColorProcessor cip = new ColorProcessor(w, h);
3 ImagePlus cimg = new ImagePlus("My New Color Image", cip);
4 cimg.show();

```

When needed, the color image can be displayed by creating an instance of the class `ImagePlus` (line 3) and calling its `show()` method. Since `cip` is of type `ColorProcessor`, the resulting `ImagePlus` object `cimg` is also a color image. The following code segment demonstrates how this could be verified:

```
5 if (cimg.getType() == ImagePlus.COLOR_RGB) {  
6     int b = cimg.getBitDepth(); // b = 24  
7     IJ.write("this is an RGB color image with " + b + " bits");  
8 }
```

Indexed color images

The structure of an indexed image in ImageJ is given in Fig. 8.5, where each element of the index array is 8 bits and therefore can represent a maximum of 256 different colors. When programming, indexed images are similar to grayscale images, as both make use of a color table to determine the actual color of the pixel. Indexed images differ from grayscale images only in that the contents of the color table are not intensity values but RGB values.

Opening and saving indexed images. ImageJ supports the following types of image formats for indexed images:

- **GIF**: index values with 1 to 8 bits (2 to 256 colors), 3×8 -bit color values.
- **PNG**: index values with 1 to 8 Bits (2 to 256 colors), 3×8 -bit color values.
When saved as PNG, indexed images are stored as full-color RGB images.
- **BMP, TIFF** (uncompressed): index values with 1 to 8 bits (2 to 256 colors), 3×8 -bit color values.

Working with indexed images. The indexed format is mostly used as a space-saving means of image storage and is not directly useful as a processing format since an index value in the pixel array is arbitrarily related to the actual color, found in the color table, that it represents. When working with indexed images it usually makes no sense to base any numerical interpretations on the pixel values or to apply any filter operations designed for 8-bit intensity images. Figure 8.8 illustrates an example of applying a Gaussian filter and a median filter to the pixels of an indexed image. Since there is no meaningful quantitative relation between the actual colors and the index values, the results are erratic. Note that even the use of the median filter is inadmissible because no ordering relation exists between the index values. Thus, with few exceptions, ImageJ functions do not permit the application of such operations to indexed images. Generally, when processing an indexed image, you first convert it into a true color RGB image and then after processing convert it back into an indexed image.

```

1 // File Brighten_Index_Image.java
2
3 import ij.ImagePlus;
4 import ij.WindowManager;
5 import ij.plugin.filter.PlugInFilter;
6 import ij.process.ImageProcessor;
7 import java.awt.image.IndexColorModel;
8
9 public class Brighten_Index_Image implements PlugInFilter {
10
11    public int setup(String arg, ImagePlus im) {
12        return DOES_8C; // this plugin works on indexed color images
13    }
14
15    public void run(ImageProcessor ip) {
16        IndexColorModel icm =
17            (IndexColorModel) ip.getColorModel();
18        int pixBits = icm.getPixelSize();
19        int mapSize = icm.getMapSize();
20
21        //retrieve the current lookup tables (maps) for R,G,B
22        byte[] Rmap = new byte[mapSize]; icm.getReds(Rmap);
23        byte[] Gmap = new byte[mapSize]; icm.getGreens(Gmap);
24        byte[] Bmap = new byte[mapSize]; icm.getBlues(Bmap);
25
26        //modify the lookup tables
27        for (int idx = 0; idx < mapSize; idx++){
28            int r = 0xff & Rmap[idx]; //mask to treat as unsigned byte
29            int g = 0xff & Gmap[idx];
30            int b = 0xff & Bmap[idx];
31            Rmap[idx] = (byte) Math.min(r + 10, 255);
32            Gmap[idx] = (byte) Math.min(g + 10, 255);
33            Bmap[idx] = (byte) Math.min(b + 10, 255);
34        }
35
36        //create a new color model and apply to the image
37        IndexColorModel icm2 =
38            new IndexColorModel(pixBits, mapSize, Rmap, Gmap, Bmap);
39        ip.setColorModel(icm2);
40
41        //update the resulting image
42        WindowManager.getCurrentImage().updateAndDraw();
43    }
44
45 } // end of class Brighten_Index_Image

```

Program 8.3 Working with indexed images (ImageJ plugin). This plugin increases the brightness of an image by 10 units by modifying the image's color table (palette). The actual values in the pixel array, which are indices into the palette, are not changed.

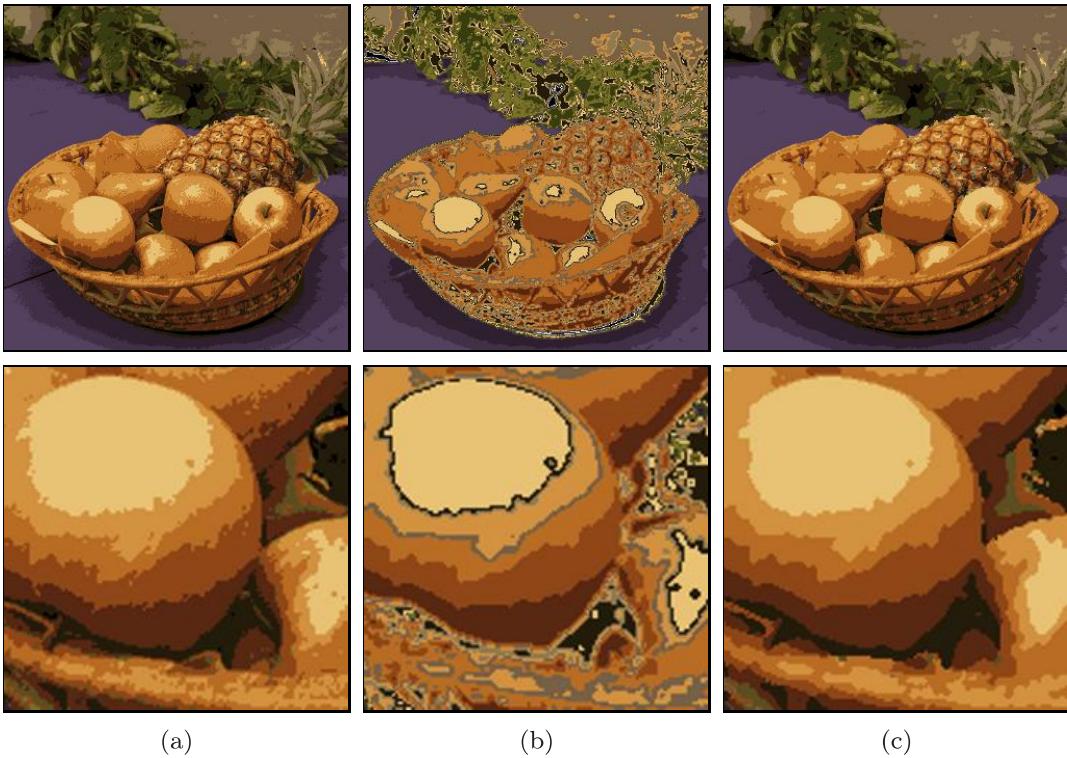


Figure 8.8 Improper application of smoothing filters to an indexed color image. Indexed image with 16 colors (a) and results of applying a linear smoothing filter (b) and a 3×3 median filter (c) to the pixel array (that is, the *index* values). The application of a linear filter makes no sense, of course, since no meaningful relation exists between the index values in the pixel array and the actual image intensities. While the median filter (c) delivers seemingly plausible results in this case, its use is also inadmissible because no suitable ordering relation exists between the index values.

When an ImageJ plugin is supposed to process indexed images, its `setup()` method should return the `DOES_8C` (“8-bit color”) flag. The plugin in Prog. 8.3 shows how to increase the intensity of the three color components of an indexed image by 10 units (analogously to Progs. 8.1 and 8.2 for RGB images). Notice how in indexed images only the palette is modified and the original pixel data, the index values, remain the same. The color table of `ImageProcessor` is accessible through a `ColorModel`⁵ object, which can be read using the method `getColorModel()` and modified using `setColorModel()`.

The `ColorModel` object for indexed images (as well as 8-bit grayscale images) is a subtype of `IndexColorModel`, which contains three color tables (*maps*) representing the red, green, and blue components as separate `byte` arrays. The size of these tables (2 to 256) can be determined by calling the method `getMapSize()`. Note that the elements of the palette should be interpreted as *unsigned* bytes with values ranging from 0 to 255. Just as with

⁵ Defined in the standard Java class `java.awt.image.ColorModel`.

grayscale pixel values, during the conversion to `int` values, these color component values must also be bitwise masked with `0xff` as shown in Prog. 8.3 (lines 28–30).

As a further example, Prog. 8.4 shows how to convert an indexed image to a true color RGB image of type `ColorProcessor`. Conversion in this direction poses no problems because the RGB component values for a particular pixel are simply taken from the corresponding color table entry, as described by Eqn. (8.3). On the other hand, conversion in the other direction requires *quantization* of the RGB color space and is as a rule more difficult and involved (see Vol. 2 [6, Sec. 5] for more details). In practice, most applications make use of existing conversion methods such as those available in ImageJ (see pp. 200–200).

Creating indexed images. In ImageJ, no special method is provided for the creation of indexed images, so in almost all cases they are generated by converting an existing image. The following method demonstrates how to directly create an indexed image if required:

```

1 ByteProcessor makeIndexColorImage(int w, int h, int nColors) {
2     // allocate red, green, blue color tables:
3     byte[] Rmap = new byte[nColors];
4     byte[] Gmap = new byte[nColors];
5     byte[] Bmap = new byte[nColors];
6     // color maps need to be filled here
7     byte[] pixels = new byte[w * h];
8     // pixel array (color indices) needs to be filled here
9     IndexColorModel cm
10    = new IndexColorModel(8, nColors, Rmap, Gmap, Bmap);
11    return new ByteProcessor(w, h, pixels, cm);
12 }
```

The parameter `nColors` defines the number of colors (and thus the size of the palette) and must be a value in the range of 2 to 256. To use the above template, you would complete it with code that filled the three `byte` arrays for the RGB components (`Rmap`, `Gmap`, `Bmap`) and the index array (`pixels`) with the appropriate values.

Transparency. Transparency is one of the reasons indexed images are often used for Web graphics. In an indexed image, it is possible to define one of the index values so that it is displayed in a transparent manner and at selected image locations the background beneath the image shows through. In Java this can be controlled when creating the image's color model (`IndexColorModel`). As an example, to make color index 2 in Prog. 8.3 transparent, lines 37–39 would need to be modified as follows:

```

1 // File Index_To_Rgb.java
2
3 import ij.ImagePlus;
4 import ij.plugin.filter.PlugInFilter;
5 import ij.process.ColorProcessor;
6 import ij.process.ImageProcessor;
7 import java.awt.image.IndexColorModel;
8
9 public class Index_To_Rgb implements PlugInFilter {
10    static final int R = 0, G = 1, B = 2;
11
12    public int setup(String arg, ImagePlus im) {
13        return DOES_8C + NO_CHANGES; //does not alter original image
14    }
15
16    public void run(ImageProcessor ip) {
17        int w = ip.getWidth();
18        int h = ip.getHeight();
19
20        //retrieve the color table (palette) for R,G,B
21        IndexColorModel icm =
22            (IndexColorModel) ip.getColorModel();
23        int mapSize = icm.getMapSize();
24        byte[] Rmap = new byte[mapSize]; icm.getReds(Rmap);
25        byte[] Gmap = new byte[mapSize]; icm.getGreens(Gmap);
26        byte[] Bmap = new byte[mapSize]; icm.getBlues(Bmap);
27
28        //create new 24-bit RGB image
29        ColorProcessor cp = new ColorProcessor(w,h);
30        int[] RGB = new int[3];
31        for (int v = 0; v < h; v++) {
32            for (int u = 0; u < w; u++) {
33                int idx = ip.getPixel(u, v);
34                RGB[R] = Rmap[idx];
35                RGB[G] = Gmap[idx];
36                RGB[B] = Bmap[idx];
37                cp.set(u, v, RGB);
38            }
39        }
40        ImagePlus cimg = new ImagePlus("RGB Image", cp);
41        cimg.show();
42    }
43
44 } // end of class Index_To_Rgb

```

Program 8.4 Converting an indexed image to a true color RGB image (ImageJ plugin).

```

1 int tIdx = 2; // index of transparent color
2 IndexColorModel icm2 = new
3     IndexColorModel(pixBits, mapSize, Rmap, Gmap, Bmap, tIdx);
4 ip.setColorModel(icm2);

```

At this time, however, ImageJ does not support the transparency property; it is not considered during display, and it is lost when the image is saved.

Color image conversion in ImageJ

In ImageJ, the following methods for converting between different types of color and grayscale image objects of type `ImagePlus` and processor objects of type `ImageProcessor` are available:

Converting images of type `ImageProcessor`. ImageJ objects of type `ImageProcessor` can be converted using the methods listed in Table 8.1. Each of these methods returns a new `ImageProcessor` object, unless the original image is already of the desired type. If this is the case, only a reference to the original image processor is returned, i. e., *no* duplication or modification occurs. The following example demonstrates the conversion from an `ByteProcessor` (grayscale) image type to an RGB color image:

```

1  ByteProcessor ip1; // a grayscale image
2  ...
3  ImageProcessor ip2 = ip1.convertToRGB();
4  // now ip2 is of type ColorProcessor, ip1 is unmodified.
5  ...

```

In this case, a new object (`ip2`) of type `ColorProcessor` is created and the original object (`ip1`) remains unchanged.

Converting images of type `ImagePlus`. ImageJ image objects of type `ImagePlus` can be converted with the help of methods from the ImageJ class `ImageConverter`, as summarized in Table 8.2. The following example demonstrates the conversion to an RGB color image:

```

1  import ij.process.ImageConverter;
2  ...
3  ImagePlus ipl;
4  ...
5  ImageConverter ic = new ImageConverter(ipl);
6  ic.convertToRGB();
7  // ipl is an RGB image now

```

Note that the method `convertToRGB()` does not return a new image object, but instead modifies the original `ImagePlus` object `ipl`.

8.2 Color Spaces and Color Conversion

The RGB color system is well-suited for use in programming, as it is simple to manipulate and maps directly to the typical display hardware. When modifying

Table 8.1 Conversion methods for images of type `ImageProcessor`. If `doScaling` is true in the first two methods, the pixel values are automatically scaled to the maximum range of the new image.

<code>ImageProcessor convertToByte(boolean doScaling)</code>	Converts to an 8-bit grayscale image (<code>ByteProcessor</code>).
<code>ImageProcessor convertToShort(boolean doScaling)</code>	Converts to a 16-bit grayscale image (<code>ShortProcessor</code>).
<code>ImageProcessor convertToFloat()</code>	Converts to a 32-bit floating-point image (<code>FloatProcessor</code>).
<code>ImageProcessor convertToRGB()</code>	Converts to a 32-bit RGB color image (<code>ColorProcessor</code>).

Table 8.2 Methods of the ImageJ class `ImageConverter` for converting `ImagePlus` objects. Note that these methods do not create any new images, but instead modify the original `ImagePlus` object `ipl` used to instantiate the `ImageConverter`.

<code>ImageConverter(ImagePlus ipl)</code>	Instantiates an <code>ImageConverter</code> object for the image <code>ipl</code> .
<code>void convertToGray8()</code>	Converts <code>ipl</code> to an 8-bit grayscale image.
<code>void convertToGray16()</code>	Converts <code>ipl</code> to a 16-bit grayscale image.
<code>void convertToGray32()</code>	Converts <code>ipl</code> to a 32-bit grayscale image (<code>float</code>).
<code>void convertToRGB()</code>	Converts <code>ipl</code> to an RGB color image.
<code>void convertRGBtoIndexedColor(int nColors)</code>	Converts the RGB true color image <code>ipl</code> to an indexed image with 8-bit index values and <code>nColors</code> colors, performing color quantization.
<code>void convertToHSB()</code>	Converts <code>ipl</code> to a color image using the HSB color space (see Sec. 8.2.3).
<code>void convertHSBtoRGB()</code>	Converts an HSB color space image <code>ipl</code> to an RGB color image.

colors within the RGB space, it is important to remember that the *metric*, or *measured distance* within this color space, does not proportionally correspond to our perception of color (e.g., doubling the value of the red component does not necessarily result in a color which appears to be twice as red). In general, in this space, modifying different color points by the same amount can cause very different changes in color. In addition, brightness changes in the RGB color space are also perceived as nonlinear.

Since any coordinate movement modifies color tone, saturation, and brightness all at once, color selection in RGB space is difficult and quite non-intuitive. Color selection is more intuitive in other color spaces, such as the HSV space (see Sec. 8.2.3), since perceptual color features, such as saturation, are represented individually and can be modified independently. Alternatives to the RGB color space are also used in applications such as the automatic separation of objects from a colored background (the *blue box* technique in television), encoding television signals for transmission, or in printing, and are thus also relevant in digital image processing.

Figure 8.9 shows the distribution of the colors from natural images in the RGB color space. The first half of this section introduces alternative color spaces and the methods of converting between them and later discusses the choices that need to be made to correctly convert a color image to grayscale. In addition to the classical color systems most widely used in programming, precise reference systems, such as the CIEXYZ color space, gain increasing importance in practical color processing.

8.2.1 Conversion to Grayscale

The conversion of an RGB color image to a grayscale image proceeds by computing the equivalent gray or *luminance* value Y for each RGB pixel. In its simplest form, Y could be computed as the average

$$Y = \text{Avg}(R, G, B) = \frac{R + G + B}{3} \quad (8.4)$$

of the three color components R , G , and B . Since we perceive both red and green as being substantially brighter than blue, the resulting image will appear to be too dark in the red and green areas and too bright in the blue ones. Therefore, a weighted sum of the color components

$$Y = \text{Lum}(R, G, B) = w_R \cdot R + w_G \cdot G + w_B \cdot B \quad (8.5)$$

is typically used to compute the equivalent luminance value. The weights most often used were originally developed for encoding analog color television signals (see Sec. 8.2.4):

$$w_R = 0.299, \quad w_G = 0.587, \quad w_B = 0.114. \quad (8.6)$$

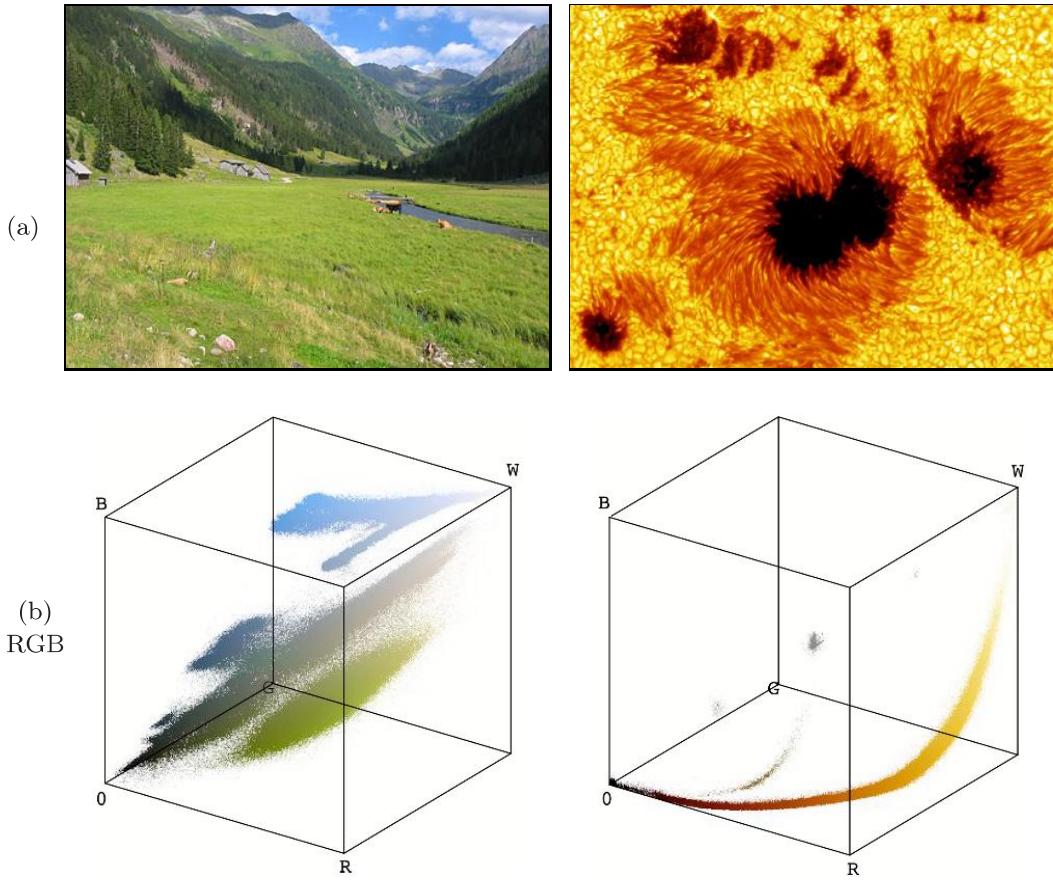


Figure 8.9 Examples of the color distribution of natural images. Original images: landscape photograph with dominant green and blue components and sun-spot image with rich red and yellow components (a). Distribution of image colors in RGB-space (b).

Those recommended in ITU-BT.709 [20] for digital color encoding are

$$w_R = 0.2125, \quad w_G = 0.7154, \quad w_B = 0.072. \quad (8.7)$$

If each color component is assigned the same weight, as in Eqn. (8.4), this is of course just a special case of Eqn. (8.5).

Note that, although these weights were developed for use with TV signals, they are optimized for *linear* RGB component values, i.e., signals with no gamma correction. In many practical situations, however, the RGB components are actually *nonlinear*, particularly when we work with sRGB images (see Vol. 2 [6, Sec. 6.3]). In this case, the RGB components must first be linearized to obtain the correct luminance values with the above weights. An alternative is to estimate the luminance without linearization by computing the weighted sum of the *nonlinear* component values and applying a different set of weights (for details see p. 109 of Vol. 2 [6, Sec. 6.3]).

In some color systems, instead of a weighted sum of the RGB color components, a nonlinear brightness function, for example the *value* V in HSV (Eqn. (8.11) in Sec. 8.2.3) or the *luminance* L in HLS (Eqn. (8.21)), is used as the intensity value Y .

Hueless (gray) color images

An RGB image is hueless or gray when the RGB components of each pixel $I(u, v) = (R, G, B)$ are the same; i. e., if

$$R = G = B.$$

Therefore, to completely remove the color from an RGB image, simply replace the R , G , and B component of each pixel with the equivalent gray value Y ,

$$I_g(u, v) \leftarrow \begin{pmatrix} R_g \\ G_g \\ B_g \end{pmatrix} = \begin{pmatrix} Y \\ Y \\ Y \end{pmatrix}, \quad (8.8)$$

by using $Y = \text{Lum}(R, G, B)$ from Eqn. (8.5), for example. The resulting grayscale image should have the same subjective brightness as the original color image.

Grayscale conversion in ImageJ

In ImageJ, the simplest way to convert an RGB color image (of type **ColorProcessor**) into an 8-bit grayscale image is to use the method

```
convertToByte(boolean doScaling),
```

which returns a new image of type **ByteProcessor** (see Table 8.1 and the example on page 200). ImageJ uses the default weights $w_R = w_G = w_B = \frac{1}{3}$ (as in Eqn. (8.4)) for the RGB components, or $w_R = 0.299$, $w_G = 0.587$, $w_B = 0.114$ (as in Eqn. (8.6)) if the “Weighted RGB Conversions” option is selected in the **Edit**→**Options**→**Conversions** dialog. Arbitrary weights (w_r , w_g , w_b) can be specified for subsequent conversion operations through the static **ColorProcessor** method

```
setWeightingFactors(double wr, double wg, double wb).
```

Similarly, the static method **ColorProcessor.getWeightingFactors()** can be used to retrieve the current weights as a 3-element **double**-array. Note that no linearization is performed on the color components, which should be considered when working with (nonlinear) sRGB colors (see Vol. 2 [6, Sec. 6.3] for details).

8.2.2 Desaturating Color Images

Desaturation is the uniform reduction of the amount of color in an RGB image in a *continuous* manner. It is done by replacing each RGB pixel by a desaturated color (R_d, G_d, B_d) computed by linear interpolation between the pixel's original color and the corresponding (Y, Y, Y) gray point in the RGB space, i. e.,

$$\begin{pmatrix} R_d \\ G_d \\ B_d \end{pmatrix} \leftarrow \begin{pmatrix} Y \\ Y \\ Y \end{pmatrix} + s_{\text{col}} \cdot \begin{pmatrix} R - Y \\ G - Y \\ B - Y \end{pmatrix}, \quad (8.9)$$

where the factor $s_{\text{col}} \in [0, 1]$ controls the remaining amount of color saturation (Fig. 8.10). A value of $s_{\text{col}} = 0$ completely eliminates all color, resulting in a true grayscale image, and with $s_{\text{col}} = 1$ the color values will be unchanged. In Prog. 8.5, continuous desaturation as defined in Eqn. (8.9) is implemented as an ImageJ plugin.

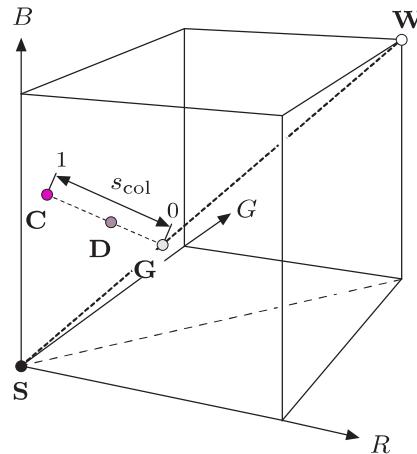


Figure 8.10 Desaturation in RGB space: original color point $\mathbf{C} = (R, G, B)$, its corresponding gray point $\mathbf{G} = (Y, Y, Y)$, and the desaturated color point $\mathbf{D} = (R_d, G_d, B_d)$. Saturation is controlled by the factor s_{col} .

8.2.3 HSV/HSB and HLS Color Space

In the **HSV** color space, colors are specified by the components *hue*, *saturation*, and *value*. Often, such as in Adobe products and the Java API, the **HSV** space is called **HSB**. While the acronym is different (in this case $B = \text{brightness}$),⁶ it denotes the same color space. The HSV color space is traditionally shown

⁶ Sometimes the HSV space is also referred to as the “HSI” space, where ‘I’ stands for *intensity*.

```

1 // File Desaturate_Rgb.java
2
3 import ij.ImagePlus;
4 import ij.plugin.filter.PlugInFilter;
5 import ij.process.ImageProcessor;
6
7 public class Desaturate_Rgb implements PlugInFilter {
8
9     static double sCol = 0.3; // color saturation factor
10
11    public int setup(String arg, ImagePlus im) {
12        return DOES_RGB;
13    }
14
15    public void run(ImageProcessor ip) {
16
17        // iterate over all pixels
18        for (int v = 0; v < ip.getHeight(); v++) {
19            for (int u = 0; u < ip.getWidth(); u++) {
20
21                // get int-packed color pixel
22                int c = ip.get(u, v);
23
24                // extract RGB components from color pixel
25                int r = (c & 0xff0000) >> 16;
26                int g = (c & 0x00ff00) >> 8;
27                int b = (c & 0x0000ff);
28
29                // compute equivalent gray value
30                double y = 0.299 * r + 0.587 * g + 0.114 * b;
31
32                // linearly interpolate (yyy) ↔ (rgb)
33                r = (int) (y + sCol * (r - y));
34                g = (int) (y + sCol * (g - y));
35                b = (int) (y + sCol * (b - y));
36
37                // reassemble color pixel
38                c = ((r & 0xff)<<16) | ((g & 0xff)<<8) | b & 0xff;
39                ip.set(u, v, c);
40            }
41        }
42    }
43
44 } // end of class Desaturate_Rgb

```

Program 8.5 Continuous desaturation of an RGB color image (ImageJ plugin). The amount of color saturation is controlled by the variable `sCol` defined in line 9 (see Eqn. (8.9)).

as an upside-down, six-sided pyramid (Fig. 8.11 (a)), where the vertical axis represents the V (brightness) value, the horizontal distance from the axis the S (saturation) value, and the angle the H (hue) value. The black point is at

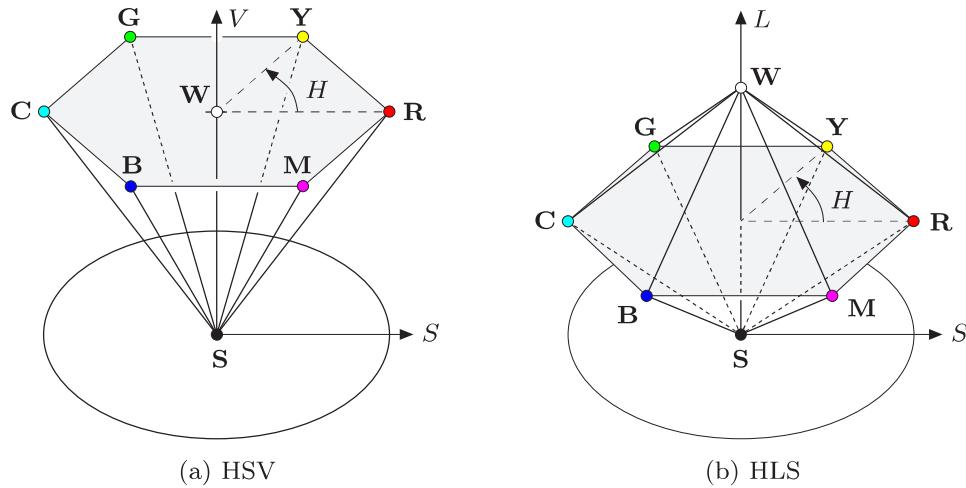


Figure 8.11 HSV and HLS color space are traditionally visualized as a single or double hexagonal pyramid. The brightness V (or L) is represented by the vertical dimension, the color saturation S by the radius from the pyramid’s axis, and the hue h by the angle. In both cases, the primary colors red (**R**), green (**G**), and blue (**B**) and the pairwise mixed colors yellow (**Y**), cyan (**C**), and magenta (**M**) lie on a common plane with black (**S**) at the tip. The essential difference between the HSV and HLS color spaces is the location of the white point (**W**).

the tip of the pyramid and the white point lies in the center of the base. The three primary colors *red*, *green*, and *blue* and the pairwise mixed colors *yellow*, *cyan* and *magenta* are the corner points of the base. While this space is often represented as a pyramid, according to its mathematical definition, the space is actually a *cylinder*, as shown below (Fig. 8.13).

The **HLS** color space⁷ (*hue*, *luminance*, *saturation*) is very similar to the HSV space, and the *hue* component is in fact completely identical in both spaces. The *luminance* and *saturation* values also correspond to the vertical axis and the radius, respectively, but are defined differently than in HSV space. The common representation of the HLS space is as a double pyramid (Fig. 8.11 (b)), with black on the bottom tip and white on the top. The primary colors lie on the corner points of the hexagonal base between the two pyramids. Even though it is often portrayed in this intuitive way, mathematically the HLS space is again a cylinder (see Fig. 8.15).

$RGB \rightarrow HSV$

To convert from RGB to the HSV color space, we first find the *saturation* of the RGB color components $R, G, B \in [0, C_{\max}]$, with C_{\max} being the maximum

⁷ The acronyms HLS and HSL are used interchangeably.

component value (typically 255), as

$$S_{\text{HSV}} = \begin{cases} \frac{C_{\text{rng}}}{C_{\text{high}}} & \text{for } C_{\text{high}} > 0 \\ 0 & \text{otherwise,} \end{cases} \quad (8.10)$$

and the luminance (*value*)

$$V_{\text{HSV}} = \frac{C_{\text{high}}}{C_{\text{max}}}, \quad (8.11)$$

with C_{high} , C_{low} , and C_{rng} defined as

$$C_{\text{high}} = \max(R, G, B), \quad C_{\text{low}} = \min(R, G, B), \quad C_{\text{rng}} = C_{\text{high}} - C_{\text{low}}. \quad (8.12)$$

Finally, we need to specify the *hue* value H_{HSV} . When all three RGB color components have the same value ($R = G = B$), then we are dealing with an *achromatic* (gray) pixel. In this particular case $C_{\text{rng}} = 0$ and thus the saturation value $S_{\text{HSV}} = 0$, consequently the hue is undefined. To compute H_{HSV} when $C_{\text{rng}} > 0$, we first normalize each component using

$$R' = \frac{C_{\text{high}} - R}{C_{\text{rng}}}, \quad G' = \frac{C_{\text{high}} - G}{C_{\text{rng}}}, \quad B' = \frac{C_{\text{high}} - B}{C_{\text{rng}}}. \quad (8.13)$$

Then, depending on which of the three original color components had the maximal value, we compute a preliminary hue H' as

$$H' = \begin{cases} B' - G' & \text{if } R = C_{\text{high}} \\ R' - B' + 2 & \text{if } G = C_{\text{high}} \\ G' - R' + 4 & \text{if } B = C_{\text{high}}. \end{cases} \quad (8.14)$$

Since the resulting value for H' lies on the interval $[-1 \dots 5]$, we obtain the final hue value by normalizing to the interval $[0, 1]$ as

$$H_{\text{HSV}} = \frac{1}{6} \cdot \begin{cases} (H' + 6) & \text{for } H' < 0 \\ H' & \text{otherwise.} \end{cases} \quad (8.15)$$

Hence all three components H_{HSV} , S_{HSV} , and V_{HSV} will lie within the interval $[0, 1]$. The hue value H_{HSV} can naturally also be computed in another angle interval, for example in the 0 to 360° interval using

$$H_{\text{HSV}}^\circ = H_{\text{HSV}} \cdot 360.$$

Under this definition, the RGB space unit cube is mapped to a *cylinder* with height and radius of length 1 (Fig. 8.13). In contrast to the traditional representation (Fig. 8.11), all HSB points within the entire cylinder correspond to valid color coordinates in RGB space. The mapping from RGB to the HSV space is nonlinear, as can be noted by examining how the black point stretches

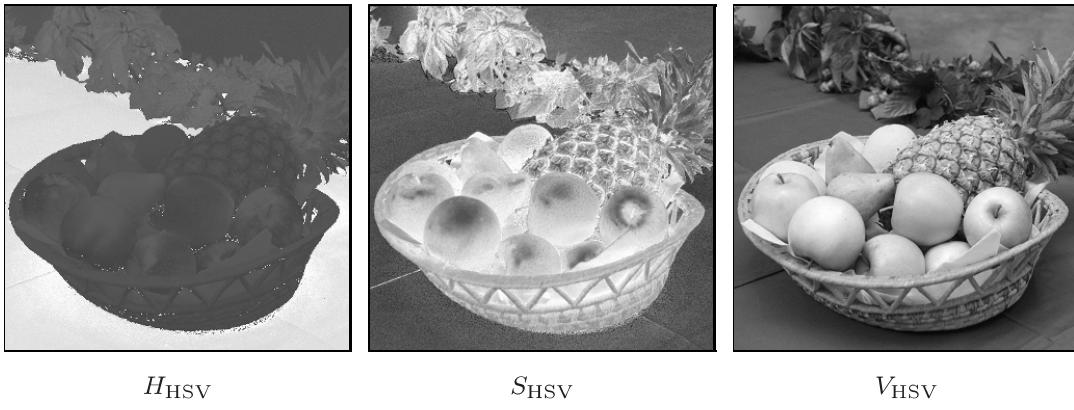


Figure 8.12 HSV components for the test image in Fig. 8.2. The darker areas in the h_{HSV} component correspond to the red and yellow colors, where the hue angle is near zero.

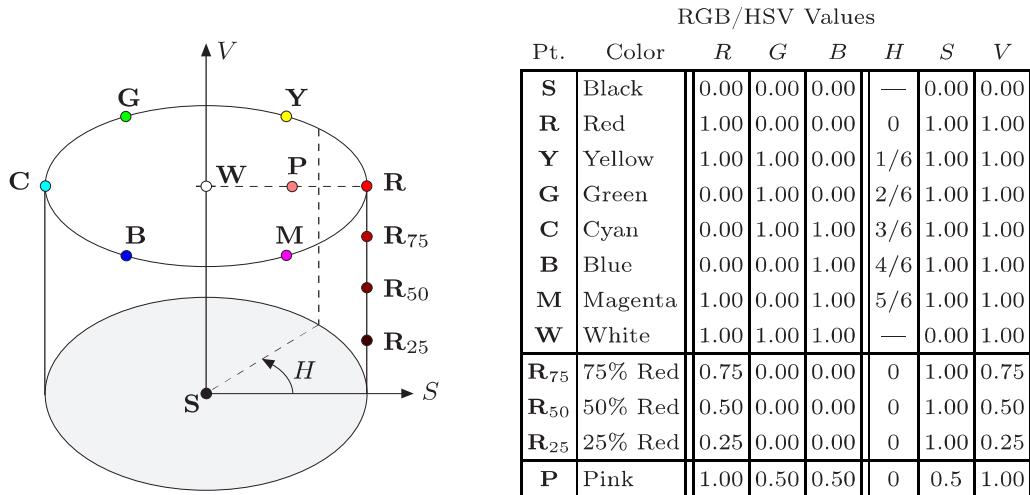


Figure 8.13 HSV color space. The illustration shows the HSV color space as a cylinder with the coordinates H (hue) as the angle, S (saturation) as the radius, and V (brightness value) as the distance along the vertical axis, which runs between the black point **S** and the white point **W**. The table lists the (R, G, B) and (H, S, V) values of the color points marked on the graphic. Pure colors (composed of only one or two components) lie on the outer wall of the cylinder ($S = 1$), as exemplified by the gradually saturated reds ($\mathbf{R}_{25}, \mathbf{R}_{50}, \mathbf{R}_{75}, \mathbf{R}$).

completely across the cylinder's base. Figure 8.12 shows the individual HSV components (in grayscale) of the test image in Fig. 8.2. Figure 8.13 plots the location of some notable color points and compares them with their locations in RGB space (see also Fig. 8.1).

Java implementation. In Java, the RGB-HSV conversion is implemented in the class `java.awt.Color` by the method

```
float[] RGBtoHSB (int r, int g, int b, float[] hsv)
```

(HSV and HSB denote the same color space). The method takes three `int` arguments `r`, `g`, `b` (within the range $[0, 255]$) and returns a `float` array with the resulting H, S, V values in the interval $[0, 1]$. When an existing `float` array is passed as the argument `hsv`, then the result is placed in it; otherwise (when `hsv = null`) a new array is created. Here is a simple example:

```

1 import java.awt.Color;
2 ...
3 float[] hsv = new float[3];
4 int red = 128, green = 255, blue = 0;
5 hsv = Color.RGBtoHSB(red, green, blue, hsv);
6 float h = hsv[0];
7 float s = hsv[1];
8 float v = hsv[2];
9 ...

```

A possible implementation of the Java method `RGBtoHSB()` using the definition in Eqns. (8.11)–(8.15) is given in Prog. 8.6.

$HSV \rightarrow RGB$

To convert an HSV tuple $(H_{HSV}, S_{HSV}, V_{HSV})$, where H_{HSV} , S_{HSV} , and $V_{HSV} \in [0, 1]$, into the corresponding (R, G, B) color values, the appropriate color sector

$$H' = (6 \cdot H_{HSV}) \bmod 6 \quad (8.16)$$

$(0 \leq H' < 6)$ is determined first, followed by computing the intermediate values

$$\begin{aligned} c_1 &= \lfloor H' \rfloor, & x &= (1 - S_{HSV}) \cdot v, \\ c_2 &= H' - c_1, & y &= (1 - (S_{HSV} \cdot c_2)) \cdot V_{HSV}, \\ & & z &= (1 - (S_{HSV} \cdot (1 - c_2))) \cdot V_{HSV}. \end{aligned} \quad (8.17)$$

Depending on the value of c_1 , the normalized RGB values $R', G', B' \in [0, 1]$ are then computed from $v = V_{HSV}$, x , y , and z as follows:⁸

$$(R', G', B') = \begin{cases} (v, z, x) & \text{if } c_1 = 0 \\ (y, v, x) & \text{if } c_1 = 1 \\ (x, v, z) & \text{if } c_1 = 2 \\ (x, y, v) & \text{if } c_1 = 3 \\ (z, x, v) & \text{if } c_1 = 4 \\ (v, x, y) & \text{if } c_1 = 5. \end{cases} \quad (8.18)$$

⁸ The variables x , y , z used here have no relation to those used in the CIEXYZ color space.

```

1  static float[] RGBtoHSV (int R, int G, int B, float[] HSV) {
2      // R, G, B ∈ [0, 255]
3      float H = 0, S = 0, V = 0;
4      float cMax = 255.0f;
5      int cHi = Math.max(R,Math.max(G,B)); // highest color value
6      int cLo = Math.min(R,Math.min(G,B)); // lowest color value
7      int cRng = cHi - cLo;           // color range
8
9      // compute value V
10     V = cHi / cMax;
11
12     // compute saturation S
13     if (cHi > 0)
14         S = (float) cRng / cHi;
15
16     // compute hue H
17     if (cRng > 0) { // hue is defined only for color pixels
18         float rr = (float)(cHi - R) / cRng;
19         float gg = (float)(cHi - G) / cRng;
20         float bb = (float)(cHi - B) / cRng;
21         float hh;
22         if (R == cHi)                      // R is highest color value
23             hh = bb - gg;
24         else if (G == cHi)                // G is highest color value
25             hh = rr - bb + 2.0f;
26         else                            // B is highest color value
27             hh = gg - rr + 4.0f;
28         if (hh < 0)
29             hh= hh + 6;
30         H = hh / 6;
31     }
32
33     if (HSV == null) // create a new HSV array if needed
34         HSV = new float[3];
35     HSV[0] = H; HSV[1] = S; HSV[2] = V;
36     return HSV;
37 }
```

Program 8.6 RGB→HSV conversion. This Java method for RGB→HSV conversion follows the process given in the text to compute a single color tuple. It takes the same arguments and returns results identical to the standard `Color.RGBtoHSB()` method.

The scaling of the RGB components to whole numbers in the range $[0, N - 1]$ (typically $N = 256$) is carried out as follows:

$$\begin{aligned} R &= \min(\text{round}(N \cdot R'), N - 1), \\ G &= \min(\text{round}(N \cdot G'), N - 1), \\ B &= \min(\text{round}(N \cdot B'), N - 1). \end{aligned} \tag{8.19}$$

```

1  static int HSVtoRGB (float h, float s, float v) {
2      // h, s, v ∈ [0, 1]
3      float rr = 0, gg = 0, bb = 0;
4      float hh = (6 * h) % 6;                      // h' ← (6 · h) mod 6
5      int c1 = (int) hh;                            // c1 ← ⌊h'⌋
6      float c2 = hh - c1;
7      float x = (1 - s) * v;
8      float y = (1 - (s * c2)) * v;
9      float z = (1 - (s * (1 - c2))) * v;
10     switch (c1) {
11         case 0: rr=v; gg=z; bb=x; break;
12         case 1: rr=y; gg=v; bb=x; break;
13         case 2: rr=x; gg=v; bb=z; break;
14         case 3: rr=x; gg=y; bb=v; break;
15         case 4: rr=z; gg=x; bb=v; break;
16         case 5: rr=v; gg=x; bb=y; break;
17     }
18     int N = 256;
19     int r = Math.min(Math.round(rr*N), N-1);
20     int g = Math.min(Math.round(gg*N), N-1);
21     int b = Math.min(Math.round(bb*N), N-1);
22     // create int-packed RGB color:
23     int rgb = ((r&0xff)<<16) | ((g&0xff)<<8) | b&0xff;
24     return rgb;
25 }
```

Program 8.7 HSV→RGB conversion. This Java method takes the same arguments and returns identical results as the standard method `Color.HSBtoRGB()`.

Java implementation. In Java, HSV→RGB conversion is implemented in the standard AWT class `java.awt.Color` by the method

```
int HSBtoRGB (float h, float s, float v) ,
```

which takes three `float` arguments $h, s, v \in [0, 1]$ and returns the corresponding RGB color as an `int` value with 3×8 bits arranged in the standard Java RGB format (see Fig. 8.6). One possible implementation of this method is shown in Prog. 8.7.

RGB→HLS

In the HLS model, the *hue* value H_{HLS} is computed in the same way as in the HSV model (Eqns. (8.13)–(8.15)), i. e.,

$$H_{\text{HLS}} = H_{\text{HSV}}. \quad (8.20)$$

The other values, L_{HLS} and S_{HLS} , are computed as follows (for C_{high} , C_{low} , and C_{rng} , see Eqn. (8.12)):

$$L_{\text{HLS}} = \frac{C_{\text{high}} + C_{\text{low}}}{2} , \quad (8.21)$$

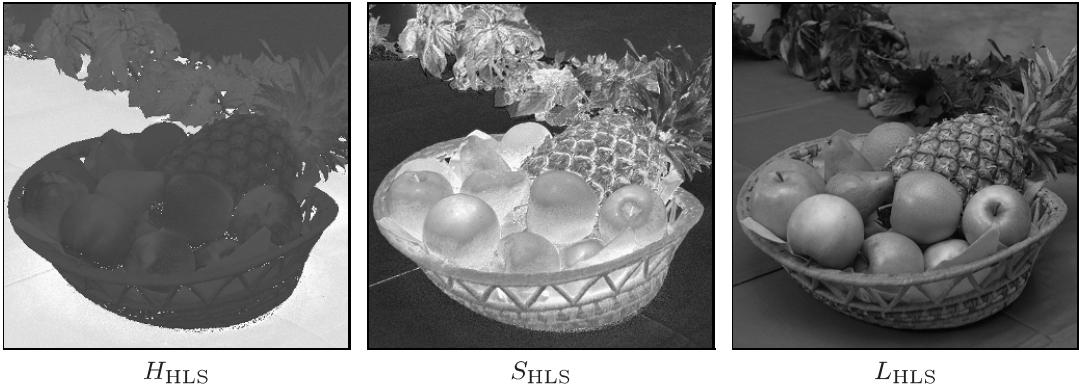


Figure 8.14 HLS color components H_{HLS} (*hue*), S_{HLS} (*saturation*), and L_{HLS} (*luminance*). Note that the S and L images are swapped to appear in the same order as in HSV space (see Fig. 8.12).

$$S_{\text{HLS}} = \begin{cases} 0 & \text{for } L_{\text{HLS}} = 0 \\ 0.5 \cdot \frac{C_{\text{rng}}}{L_{\text{HLS}}} & \text{for } 0 < L_{\text{HLS}} \leq 0.5 \\ 0.5 \cdot \frac{C_{\text{rng}}}{1-L_{\text{HLS}}} & \text{for } 0.5 < L_{\text{HLS}} < 1 \\ 0 & \text{for } L_{\text{HLS}} = 1. \end{cases} \quad (8.22)$$

Figure 8.14 shows the individual HLS components of the test image as grayscale images. Using the above definitions, the unit cube in the RGB space is again mapped to a cylinder with height and length 1 (Fig. 8.15). In contrast to the HSV space (Fig. 8.13), the primary colors lie together in the horizontal plane at $L_{\text{HLS}} = 0.5$ and the white point lies outside of this plane at $L_{\text{HLS}} = 1.0$. Using these nonlinear transformations, the black and the white points are mapped to the top and the bottom planes of the cylinder, respectively.

HLS→*RGB*

When converting from HLS to the RGB space, we assume that $H_{\text{HLS}}, S_{\text{HLS}}, L_{\text{HLS}} \in [0, 1]$. In the case where $L_{\text{HLS}} = 0$ or $L_{\text{HLS}} = 1$, the result is

$$(R', G', B') = \begin{cases} (0, 0, 0) & \text{for } L_{\text{HLS}} = 0 \\ (1, 1, 1) & \text{for } L_{\text{HLS}} = 1. \end{cases} \quad (8.23)$$

Otherwise, we again determine the appropriate color sector

$$H' = (6 \cdot H_{\text{HLS}}) \bmod 6, \quad (8.24)$$

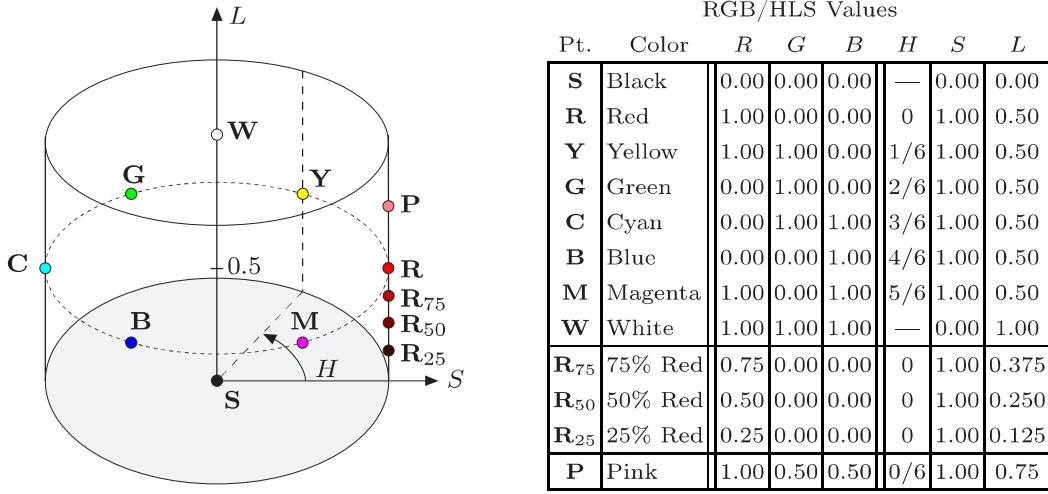


Figure 8.15 HLS color space. The illustration shows the HLS color space visualized as a cylinder with the coordinates H (hue) as the angle, S (saturation) as the radius, and L (lightness) as the distance along the vertical axis, which runs between the black point \mathbf{S} and the white point \mathbf{W} . The table lists the (R, G, B) and (H, S, L) values where “pure” colors (created using only one or two color components) lie on the lower half of the outer cylinder wall ($S = 1$), as illustrated by the gradually saturated reds (\mathbf{R}_{25} , \mathbf{R}_{50} , \mathbf{R}_{75} , \mathbf{R}). Mixtures of all three primary colors, where at least one of the components is completely saturated, lie along the upper half of the outer cylinder wall; for example, the point \mathbf{P} (pink).

where $(0 \leq H' < 6)$, and then, based on the resulting sector, we determine the values

$$\begin{aligned} c_1 &= \lfloor H' \rfloor & d &= \begin{cases} S_{\text{HLS}} \cdot L_{\text{HLS}} & \text{for } L_{\text{HLS}} \leq 0.5 \\ S_{\text{HLS}} \cdot (L_{\text{HLS}} - 1) & \text{for } L_{\text{HLS}} > 0.5 \end{cases} \\ c_2 &= H' - c_1 & w &= L_{\text{HLS}} + d \\ w &= L_{\text{HLS}} + d & y &= w - (w - x) \cdot c_2 \\ x &= L_{\text{HLS}} - d & z &= x + (w - x) \cdot c_2. \end{aligned} \quad (8.25)$$

The assignment of the RGB values is done similarly to Eqn. (8.18), i.e.,

$$(R', G', B') = \begin{cases} (w, z, x) & \text{if } c_1 = 0 \\ (y, w, x) & \text{if } c_1 = 1 \\ (x, w, z) & \text{if } c_1 = 2 \\ (x, y, w) & \text{if } c_1 = 3 \\ (z, x, w) & \text{if } c_1 = 4 \\ (w, x, y) & \text{if } c_1 = 5. \end{cases} \quad (8.26)$$

Finally, scaling the normalized $([0, 1]) R', G', B'$ color components back into the $[0, 255]$ range is done as in Eqn. (8.19).

```

1  static float[] RGBtoHLS (float R, float G, float B) {
2      // R, G, B assumed to be in [0, 1]
3      float cHi = Math.max(R,Math.max(G,B)); // highest color value
4      float cLo = Math.min(R,Math.min(G,B)); // lowest color value
5      float cRng = cHi - cLo;           // color range
6
7      // compute luminance L
8      float L = (cHi + cLo)/2;
9
10     // compute saturation S
11     float S = 0;
12     if (0 < L && L < 1) {
13         float d = (L <= 0.5f) ? L : (1 - L);
14         S = 0.5f * cRng / d;
15     }
16
17     // compute hue H
18     float H=0;
19     if (cHi > 0 && cRng > 0) {    // a color pixel
20         float rr = (float)(cHi - R) / cRng;
21         float gg = (float)(cHi - G) / cRng;
22         float bb = (float)(cHi - B) / cRng;
23         float hh;
24         if (R == cHi)                  // R is highest color value
25             hh = bb - gg;
26         else if (G == cHi)            // G is highest color value
27             hh = rr - bb + 2.0f;
28         else                          // B is highest color value
29             hh = gg - rr + 4.0f;
30
31         if (hh < 0)
32             hh= hh + 6;
33         H = hh / 6;
34     }
35
36     return new float[] {H,L,S};
37 }
```

Program 8.8 RGB→HLS conversion (Java method).

Java implementation ($RGB \leftrightarrow HLS$)

Currently there is no method in either the standard Java API or ImageJ for converting color values between RGB and HLS. Program 8.8 gives one possible implementation of the RGB→HLS conversion that follows the definitions in Eqns. (8.20)–(8.22). The HLS→RGB conversion is given in Prog. 8.9.

```

1  static float[] HLStoRGB (float H, float L, float S) {
2      // H, L, S assumed to be in [0, 1]
3      float R = 0, G = 0, B = 0;
4
5      if (L <= 0)          // black
6          R = G = B = 0;
7      else if (L >= 1)    // white
8          R = G = B = 1;
9      else {
10         float hh = (6 * H) % 6;
11         int c1 = (int) hh;
12         float c2 = hh - c1;
13         float d = (L <= 0.5f) ? (S * L) : (S * (1 - L));
14         float w = L + d;
15         float x = L - d;
16         float y = w - (w - x) * c2;
17         float z = x + (w - x) * c2;
18         switch (c1) {
19             case 0: R=w; G=z; B=x; break;
20             case 1: R=y; G=w; B=x; break;
21             case 2: R=x; G=w; B=z; break;
22             case 3: R=x; G=y; B=w; break;
23             case 4: R=z; G=x; B=w; break;
24             case 5: R=w; G=x; B=y; break;
25         }
26     }
27     return new float[] {R,G,B};
28 }
```

Program 8.9 HLS→RGB conversion (Java method).

Comparing HSV and HLS

Despite the gross similarity between the two color spaces, as Fig. 8.16 illustrates, substantial differences in the V/L and S components do exist. The essential difference between the HSV and HLS spaces is the ordering of the colors that lie between the white point **W** and the “pure” colors (**R**, **G**, **B**, **Y**, **C**, **M**), which consist of at most two primary colors, at least one of which is completely saturated.

The difference in how colors are distributed in RGB, HSV, and HLS space is readily apparent in Fig. 8.17. The starting point was a distribution of 1331 ($11 \times 11 \times 11$) color tuples obtained by uniformly sampling the RGB space at an interval of 0.1 in each dimension. The color distributions in HSV-space for a set of natural images are shown in Fig. 8.18 (p. 220).

Both the HSV and HLS color spaces are widely used in practice; for instance, for selecting colors in image editing and graphics design applications. In digital image processing, they are also used for *color keying* (that is, isolating objects according to their *hue*) on a homogeneously colored background where the

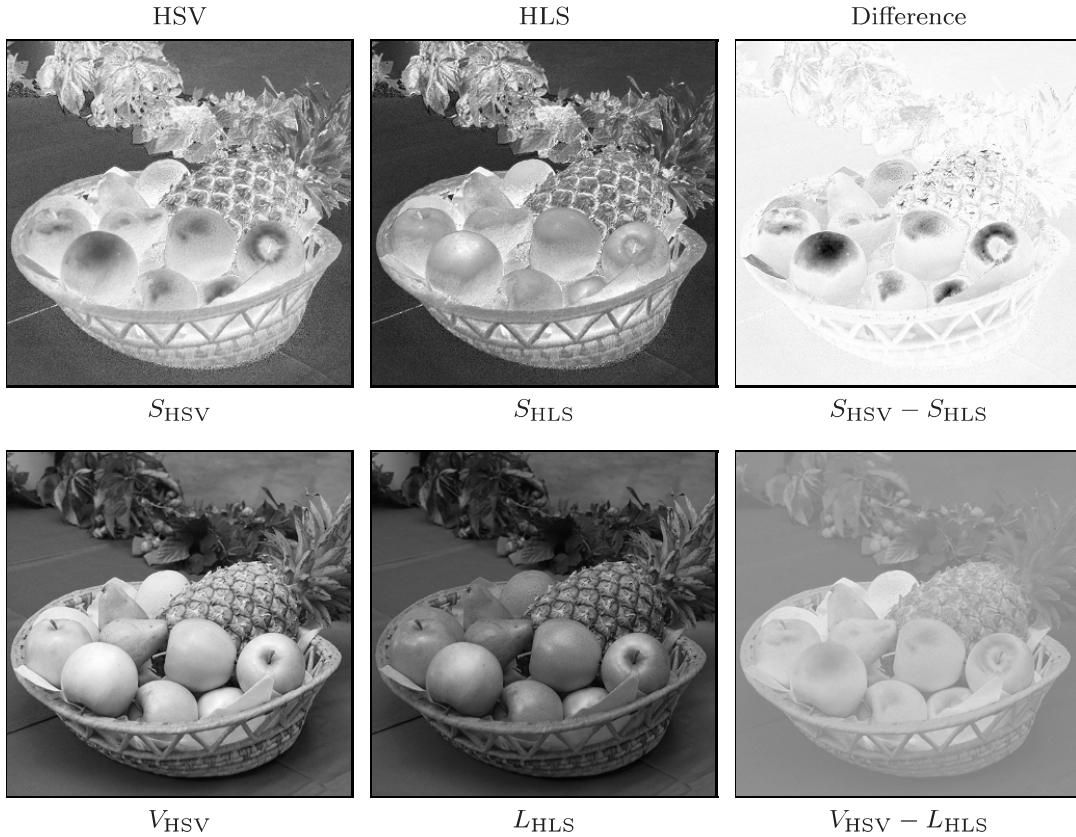


Figure 8.16 Comparison between HSV and HLS components: *saturation* (top row) and *intensity* (bottom row). In the color *saturation* difference image $S_{\text{HSV}} - S_{\text{HLS}}$ (top), light areas correspond to positive values and dark areas to negative values. Saturation in the HLS representation, especially in the brightest sections of the image, is notably higher, resulting in negative values in the difference image. For the *intensity* (*value* and *luminance*, respectively) in general, $V_{\text{HSV}} \geq L_{\text{HLS}}$ and therefore the difference $V_{\text{HSV}} - L_{\text{HLS}}$ (bottom) is always positive. The *hue* component H (not shown) is identical in both representations.

brightness is not necessarily constant.

8.2.4 TV Color Spaces—YUV, YIQ, and YC_bC_r

These color spaces are an integral part of the standards surrounding the recording, storage, transmission, and display of television signals. YUV and YIQ are the fundamental color-encoding methods for the analog NTSC and PAL systems, and YC_bC_r is a part of the international standards governing digital television [19]. All of these color spaces have in common the idea of separating the luminance component Y from two chroma components and, instead of directly encoding colors, encoding color differences. In this way, compatibility with legacy black and white systems is maintained while at the same time the bandwidth of the signal can be optimized by using different transmission bandwidths for the brightness and the color components. Since the human

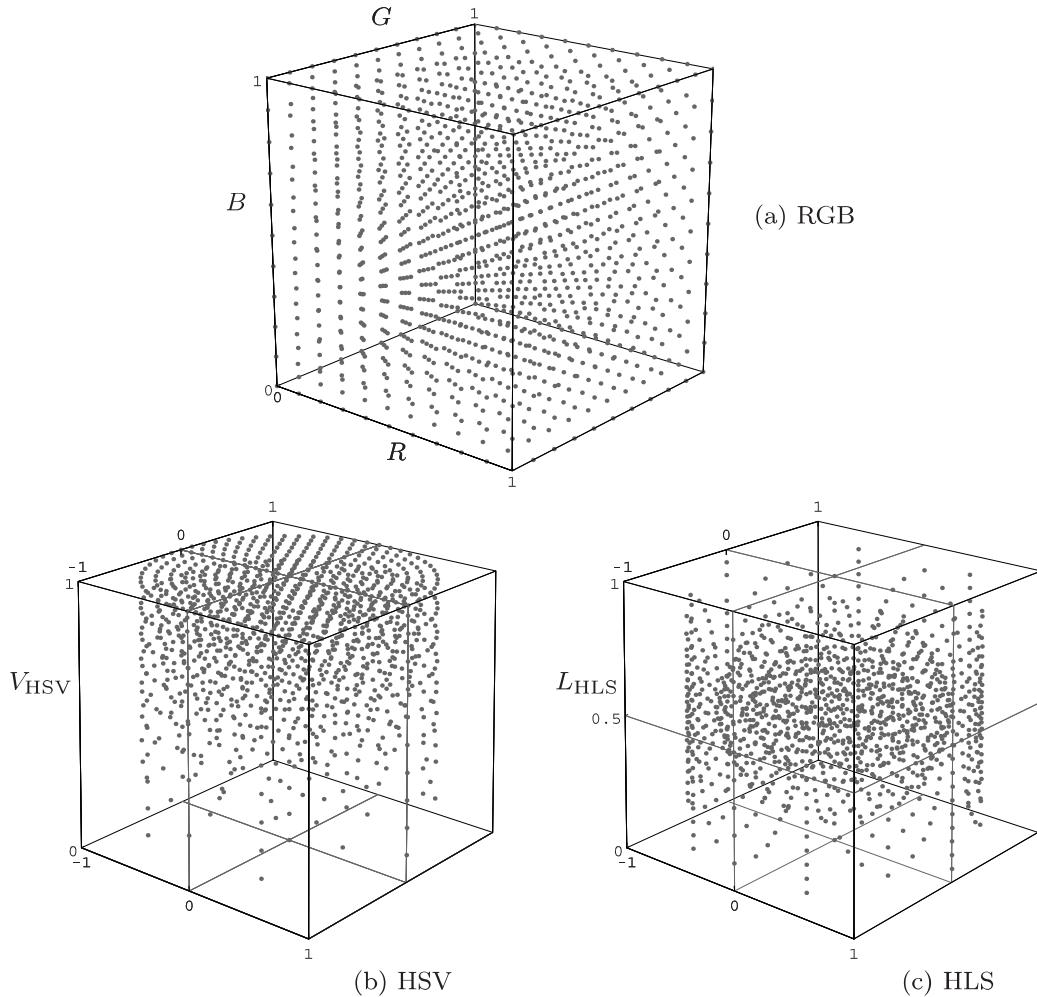


Figure 8.17 Distribution of colors in the RGB, HSV, and HLS spaces. The starting point is the uniform distribution of colors in RGB space (a). The corresponding colors in the cylindrical spaces are distributed nonsymmetrically in HSV (b) and symmetrically in HLS (c).

visual system is not able to perceive detail in the color components as well as it does in the intensity part of a video signal, the amount of information, and consequently bandwidth, used in the color channel can be reduced to approximately 1/4 of that used for the intensity component. This fact is also used when compressing digital still images and is why, for example, the JPEG codec converts RGB images to YC_bC_r. That is why these color spaces are important in digital image processing, even though raw YIQ or YUV images are rarely encountered in practice.

YUV

YUV is the basis for the color encoding used in analog television in both the North American NTSC and the European PAL systems. The luminance component Y is computed, just as in Eqn. (8.6), from the RGB components as

$$Y = 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B \quad (8.27)$$

under the assumption that the RGB values have already been gamma corrected according to the TV encoding standard ($\gamma_{\text{NTSC}} = 2.2$ and $\gamma_{\text{PAL}} = 2.8$, see Sec. 4.7) for playback. The UV components are computed from a weighted difference between the luminance and the blue or red components as

$$U = 0.492 \cdot (B - Y) \quad \text{and} \quad V = 0.877 \cdot (R - Y), \quad (8.28)$$

and the entire transformation from RGB to YUV is

$$\begin{pmatrix} Y \\ U \\ V \end{pmatrix} = \begin{pmatrix} 0.299 & 0.587 & 0.114 \\ -0.147 & -0.289 & 0.436 \\ 0.615 & -0.515 & -0.100 \end{pmatrix} \cdot \begin{pmatrix} R \\ G \\ B \end{pmatrix}. \quad (8.29)$$

The transformation from YUV back to RGB is found by inverting the matrix in Eqn. (8.29):

$$\begin{pmatrix} R \\ G \\ B \end{pmatrix} = \begin{pmatrix} 1.000 & 0.000 & 1.140 \\ 1.000 & -0.395 & -0.581 \\ 1.000 & 2.032 & 0.000 \end{pmatrix} \cdot \begin{pmatrix} Y \\ U \\ V \end{pmatrix}. \quad (8.30)$$

The color distributions in YUV-space for a set of natural images are shown in Fig. 8.18.

YIQ

The original NTSC system used a variant of YUV called YIQ (I for “in-phase”, Q for “quadrature”), where both the U and V color vectors were rotated and mirrored such that

$$\begin{pmatrix} I \\ Q \end{pmatrix} = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix} \cdot \begin{pmatrix} \cos \beta & \sin \beta \\ -\sin \beta & \cos \beta \end{pmatrix} \cdot \begin{pmatrix} U \\ V \end{pmatrix}, \quad (8.31)$$

where $\beta = 0.576$ (33°). The Y component is the same as in YUV. Although the YIQ has certain advantages with respect to bandwidth requirements it has been completely replaced by YUV [22, p. 240].

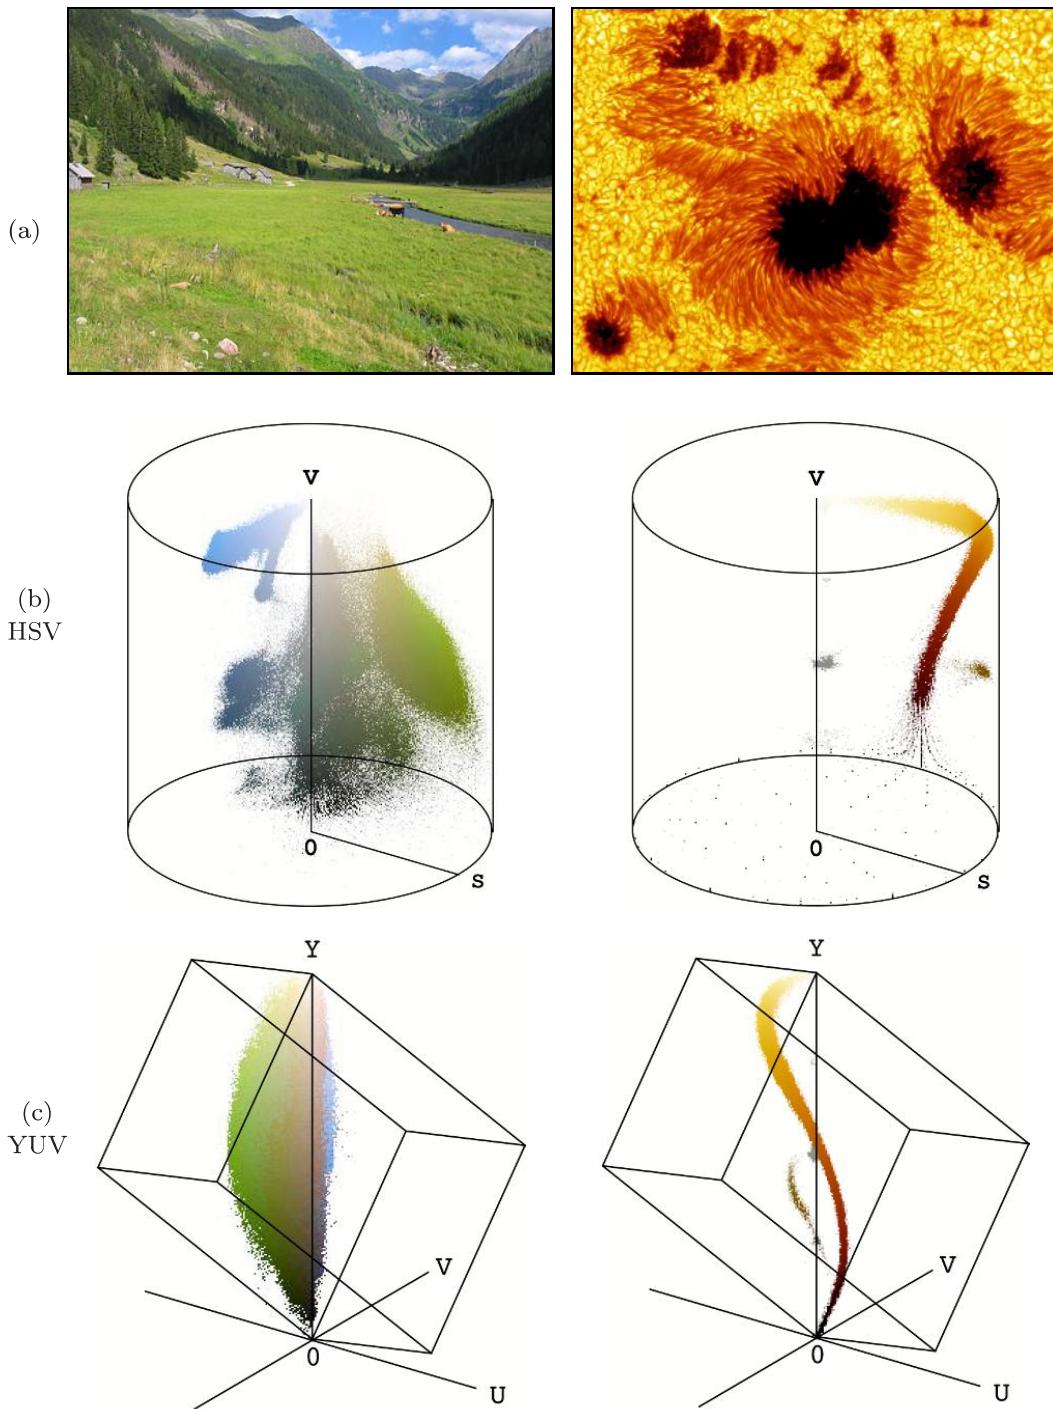


Figure 8.18 Examples of the color distribution of natural images in different color spaces. Original images (a); color distribution in HSV- (b), and YUV-space (c). See Fig. 8.9 for the corresponding distributions in RGB color space.

YC_bC_r

The YC_bC_r color space is an internationally standardized variant of YUV that is used for both digital television and image compression (for example, in JPEG). The chroma components C_b, C_r are (similar to U, V) difference values between the luminance and the blue and red components, respectively. In contrast to YUV, the weights of the RGB components for the luminance Y depend explicitly on the coefficients used for the chroma values C_b and C_r [35, p. 16]. For arbitrary weights w_B, w_R, the transformation is defined as

$$\begin{aligned} Y &= w_R \cdot R + (1 - w_B - w_R) \cdot G + w_B \cdot B, \\ C_b &= \frac{0.5}{1 - w_B} \cdot (B - Y), \\ C_r &= \frac{0.5}{1 - w_R} \cdot (R - Y), \end{aligned} \quad (8.32)$$

and the inverse transformation from YC_bC_r to RGB is

$$\begin{aligned} R &= Y + \frac{1 - w_R}{0.5} \cdot C_r, \\ G &= Y - \frac{w_B \cdot (1 - w_B) \cdot C_b - w_R \cdot (1 - w_R) \cdot C_r}{0.5 \cdot (1 - w_B - w_R)}, \\ B &= Y + \frac{1 - w_B}{0.5} \cdot C_b. \end{aligned} \quad (8.33)$$

The ITU⁹ recommendation BT.601 [21] specifies the values w_R = 0.299 and w_B = 0.114 (w_G = 1 - w_B - w_R = 0.587). Using these values, the transformation becomes

$$\begin{pmatrix} Y \\ C_b \\ C_r \end{pmatrix} = \begin{pmatrix} 0.299 & 0.587 & 0.114 \\ -0.169 & -0.331 & 0.500 \\ 0.500 & -0.419 & -0.081 \end{pmatrix} \cdot \begin{pmatrix} R \\ G \\ B \end{pmatrix}, \quad (8.34)$$

and the inverse transformation becomes

$$\begin{pmatrix} R \\ G \\ B \end{pmatrix} = \begin{pmatrix} 1.000 & 0.000 & 1.403 \\ 1.000 & -0.344 & -0.714 \\ 1.000 & 1.773 & 0.000 \end{pmatrix} \cdot \begin{pmatrix} Y \\ C_b \\ C_r \end{pmatrix}. \quad (8.35)$$

Different weights are recommended based on how the color space is used; for example, ITU-BT.709 [20] recommends w_R = 0.2125 and w_B = 0.0721 to be used in digital HDTV production. The values of U, V, I, Q, and C_b, C_r may be both positive or negative. To encode C_b, C_r values to digital numbers, a suitable offset is typically added to obtain positive-only values, e.g., 128 = 2⁷ in case of 8-bit components.

⁹ International Telecommunication Union (www.itu.int).

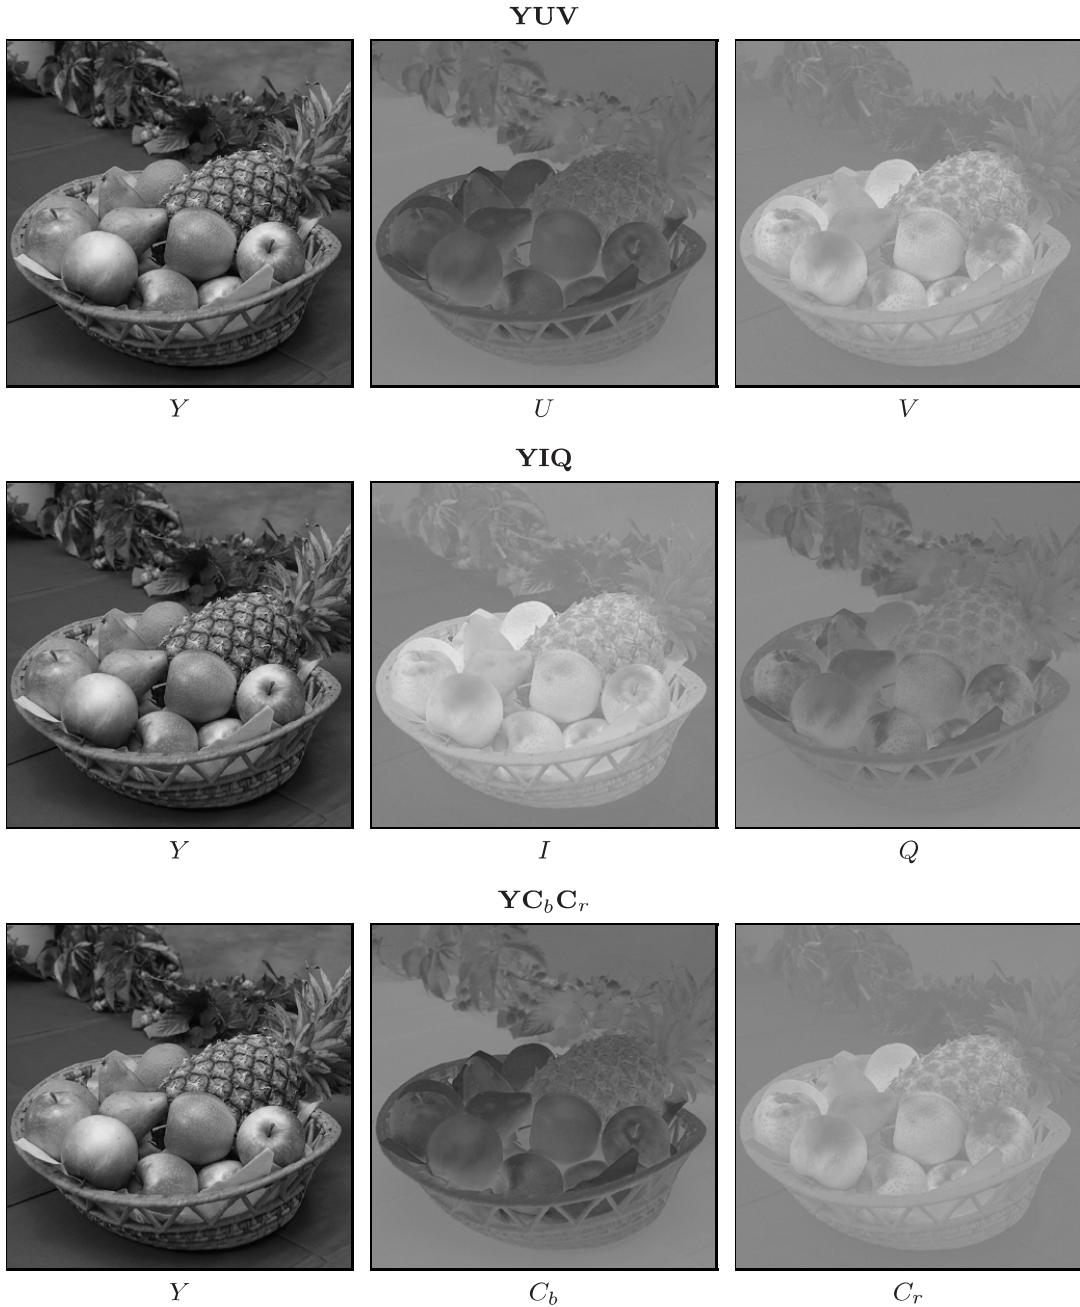


Figure 8.19 Comparing YUV-, YIQ- and $\mathbf{YC}_b\mathbf{C}_r$ -values. The Y values are identical in all three color spaces.

Figure 8.19 shows the three color spaces YUV, YIQ, and $\mathbf{YC}_b\mathbf{C}_r$ together for comparison. The U, V, I, Q , and C_b, C_r values in the right two frames have been offset by 128 so that the negative values are visible. Thus a value of zero is represented as medium gray in these images. The $\mathbf{YC}_b\mathbf{C}_r$ encoding is practically indistinguishable from YUV in these images since they both use very similar weights for the color components.

8.2.5 Color Spaces for Printing—CMY and CMYK

In contrast to the *additive* RGB color scheme (and its various color models), color printing makes use of a *subtractive* color scheme, where each printed color reduces the intensity of the reflected light at that location. Color printing requires a minimum of three primary colors; traditionally *cyan* (C), *magenta* (M) and *yellow* (Y)¹⁰ have been used.

Using subtractive color mixing on a white background, $C = M = Y = 0$ (no ink) results in the color *white* and $C = M = Y = 1$ (complete saturation of all three inks) in the color *black*. A cyan-colored ink will absorb *red* (R) most strongly, magenta absorbs *green* (G), and yellow absorbs *blue* (B). The simplest form of the CMY model is defined as

$$\begin{aligned} C &= 1 - R, \\ M &= 1 - G, \\ Y &= 1 - B. \end{aligned} \tag{8.36}$$

In practice, the color produced by fully saturating all three inks is not physically a true black. Therefore, the three primary colors C, M, Y are usually supplemented with a black ink (K) to increase the color range and coverage (gamut). In the simplest case, the amount of black is

$$K = \min(C, M, Y). \tag{8.37}$$

With rising levels of black, however, the intensity of the C, M, Y components can be gradually reduced. Many methods for reducing the primary dyes have been proposed and we look at three of them in the following.

CMY→CMYK (Version 1): In this simple variant the C, M, Y values are reduced linearly with increasing K and the modified components C', M', Y', K' are defined as

$$\begin{pmatrix} C' \\ M' \\ Y' \\ K' \end{pmatrix} = \begin{pmatrix} C-K \\ M-K \\ Y-K \\ K \end{pmatrix}. \tag{8.38}$$

CMY→CMYK (Version 2): The second variant corrects the color by reducing the C, M, Y components by $s = \frac{1}{1-K}$, resulting in stronger colors in the

¹⁰ Note that in this case Y stands for *yellow* and has nothing to do with the Y luminance component in YUV or YC_bC_r .

dark areas of the image:

$$\begin{pmatrix} C' \\ M' \\ Y' \\ K' \end{pmatrix} = \begin{pmatrix} (C-K) \cdot s \\ (M-K) \cdot s \\ (Y-K) \cdot s \\ K \end{pmatrix}, \quad \text{with } s = \begin{cases} \frac{1}{1-K} & \text{for } K < 1 \\ 1 & \text{otherwise.} \end{cases} \quad (8.39)$$

In both versions, the K -component (as defined in Eqn. (8.37)) is used directly without modification, and all gray tones (that is, when $R = G = B$) are printed using black ink K' , without any contribution from C' , M' , or Y' .

While both of these simple definitions are widely used, neither one produces high quality results. Figure 8.20 (a) compares the result from version 2 with that produced with Adobe Photoshop (Fig. 8.20 (c)). The difference in the cyan component C is particularly noticeable and also the amount of black (K) brighter areas of the image.

In practice, the required amounts of black K and C, M, Y depend so strongly on the printing process and the type of paper used that print jobs are routinely calibrated individually.

CMY→CMYK (Version 3): In print production, special transfer functions are applied to tune the results. For example, the Adobe PostScript interpreter [26, p. 345] specifies an *undercolor-removal function* $f_{UCR}(K)$ for gradually reducing the CMY components and a separate *black-generation function* $f_{BG}(K)$ for controlling the amount of black. These functions are used in the form

$$\begin{pmatrix} C' \\ M' \\ Y' \\ K' \end{pmatrix} = \begin{pmatrix} C - f_{UCR}(K) \\ M - f_{UCR}(K) \\ Y - f_{UCR}(K) \\ f_{BG}(K) \end{pmatrix}, \quad (8.40)$$

where $K = \min(C, M, Y)$ again (as defined in Eqn. (8.37)). The functions f_{UCR} and f_{BG} are usually nonlinear, and the resulting values C', M', Y', K' are scaled (typically by means of *clamping*) to the interval $[0, 1]$. The example shown in Fig. 8.20 (b) was produced using the functions

$$f_{UCR}(K) = s_K \cdot K, \quad (8.41)$$

$$f_{BG}(K) = \begin{cases} 0 & \text{for } K < K_0 \\ K_{\max} \cdot \frac{K - K_0}{1 - K_0} & \text{for } K \geq K_0, \end{cases} \quad (8.42)$$

where $s_K = 0.1$, $K_0 = 0.3$, and $K_{\max} = 0.9$ (see Fig. 8.21). With this definition, f_{UCR} reduces the CMY components by 10% of the K value (by Eqn. (8.40)), which mostly affects the dark areas of the image with high K values. The effect of the function f_{BG} (Eqn. (8.42)) is that for values of $K < K_0$ (that is in the

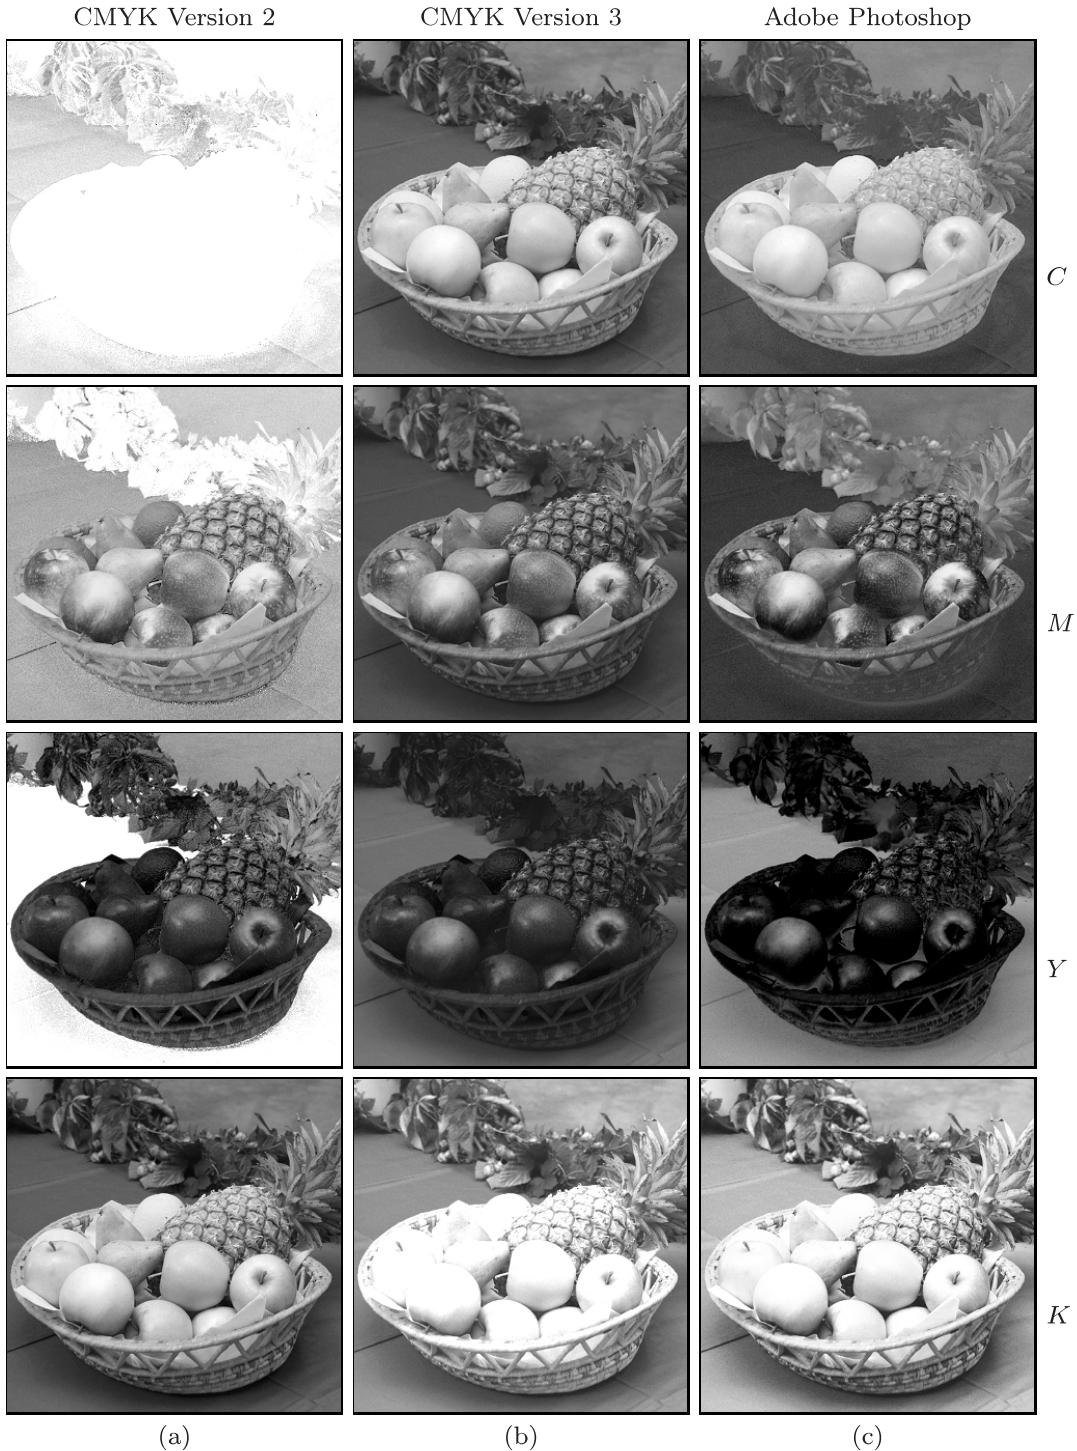


Figure 8.20 $\text{RGB} \rightarrow \text{CMYK}$ conversion comparison. Simple conversion using Eqn. (8.39) (a), applying the *undercolor-removal* and *black-generation* functions of Eqn. (8.40) (b), and results obtained with Adobe Photoshop (c). The color intensities are shown inverted, i.e., darker areas represent higher CMYK color values. The simple conversion (a), in comparison with Photoshop's result (c), shows strong deviations in all color components, C and K in particular. The results in (b) are close to Photoshop's and could be further improved by tuning the corresponding function parameters.

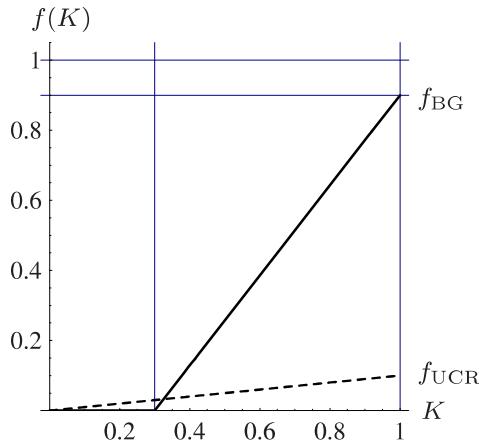


Figure 8.21 Examples of *undercolor-removal function* f_{UCR} (Eqn. (8.41)) and *black generation function* f_{BG} (Eqn. (8.42)). The parameter settings are $s_K = 0.1$, $K_0 = 0.3$, and $K_{\max} = 0.9$.

light areas of the image), no black ink is added at all. In the interval $K = K_0 \dots 1.0$, the black component is increased linearly up to the maximum value K_{\max} . The result in Fig. 8.20 (b) is relatively close to the CMYK component values produced by Photoshop¹¹ in Fig. 8.20 (c). It could be further improved by adjusting the function parameters s_K , K_0 , and K_{\max} (Eqn. (8.40)).

Even though the results of this last variant (3) for converting RGB to CMYK are better, it is only a gross approximation and still too imprecise for professional work. As we discuss in Vol. 2 [6, Sec. 6], technically correct color conversions need to be based on precise, “colorimetric” grounds.

8.3 Statistics of Color Images

8.3.1 How Many Colors Are in an Image?

A minor but frequent task in the context of color images is to determine how many different colors are contained in a given image. One way of doing this would be to create and fill a histogram array with one integer element for each color and subsequently count all histogram cells with values greater than zero. But since a 24-bit RGB color image potentially contains $2^{24} = 16,777,216$ colors, the resulting histogram array (with a size of 64 megabytes) would be larger than the image itself in most cases!

A simple solution to this problem is to *sort* the pixel values in the (one-dimensional) pixel array such that all identical colors are placed next to each

¹¹ Actually Adobe Photoshop does not convert directly from RGB to CMYK. Instead, it first converts to, and then from, the CIE L*a*b* color space (see Vol. 2 [6, Sec. 6.2]).

```

1 import ij.process.ColorProcessor;
2 import java.util.Arrays;
3
4 public class ColorStatistics {
5
6     static int countColors (ColorProcessor cp) {
7         // duplicate pixel array and sort
8         int[] pixels = ((int[]) cp.getPixels()).clone();
9         Arrays.sort(pixels);
10
11        int k = 1; // image contains at least one color
12        for (int i = 0; i < pixels.length-1; i++) {
13            if (pixels[i] != pixels[i+1])
14                k = k + 1;
15        }
16        return k;
17    }
18
19 } // end of class ColorStatistics

```

Program 8.10 Counting the colors contained in an RGB image. The method `countColors()` first creates a copy of the one-dimensional RGB (`int`) pixel array (line 8), then sorts that array, and finally counts the transitions between contiguous blocks of identical colors.

other. The sorting order is of course completely irrelevant, and the number of contiguous color blocks in the sorted pixel vector corresponds to the number of different colors in the image. This number can be obtained by simply counting the transitions between neighboring color blocks, as shown in Prog. 8.10. Of course, we do not want to sort the original pixel array (which would destroy the image) but a copy of it, which can be obtained with Java's `clone()` method.¹² Sorting of the one-dimensional array in Prog. 8.10 is accomplished (in line 9) with the generic Java method `Arrays.sort()`, which is implemented very efficiently.

8.3.2 Color Histograms

We briefly touched on histograms of color images in Sec. 3.5, where we only considered the one-dimensional distributions of the image intensity and the individual color channels. For instance, the built-in ImageJ method `getHistogram()`, when applied to an object of type `ColorProcessor`, simply computes the intensity histogram of the corresponding gray values:

```

ColorProcessor cp;
int[] H = cp.getHistogram();

```

¹² Java arrays implement the methods of the root class `Object`, including the `clone()` method specified by the `Cloneable` interface (see also Appendix B.2.5).

As an alternative, one could compute the individual intensity histograms of the three color channels, although (as discussed in Sec. 3.5.2) these do not provide any information about the actual colors in this image. Similarly, of course, one could compute the distributions of the individual components of any other color space, such as HSV or L*a*b*.

A *full* histogram of an RGB image is three-dimensional and, as noted earlier, consists of $256 \times 256 \times 256 = 2^{24}$ cells of type `int` (for 8-bit color components). Such a histogram is not only very large¹³ but also difficult to visualize.

2D color histograms

A useful alternative to the full 3D RGB histogram are two-dimensional histogram projections (Fig. 8.22). Depending on the axis of projection, we obtain 2D histograms with coordinates red-green (H_{RG}), red-blue (H_{RB}), or green-blue (H_{GB}), respectively, with the values

$$\begin{aligned} H_{RG}(r, g) &\leftarrow \text{number of pixels with } I_{RGB}(u, v) = (r, g, *), \\ H_{RB}(r, b) &\leftarrow \text{number of pixels with } I_{RGB}(u, v) = (r, *, b), \\ H_{GB}(g, b) &\leftarrow \text{number of pixels with } I_{RGB}(u, v) = (*, g, b), \end{aligned} \quad (8.43)$$

where $*$ denotes an arbitrary component value. The result is, independent of the original image size, a set of two-dimensional histograms of size 256×256 (for 8-bit RGB components), which can easily be visualized as images. Note that it is not necessary to obtain the full RGB histogram in order to compute the combined 2D histograms (see Prog. 8.11).

As the examples in Fig. 8.23 show, the combined color histograms do, to a certain extent, express the color characteristics of an image. They are therefore useful, for example, to identify the coarse type of the depicted scene or to estimate the similarity between images (see also Exercise 8.6).

8.4 Exercises

Exercise 8.1

Create an ImageJ plugin that rotates the individual components of an RGB color image; i. e., $R \rightarrow G \rightarrow B \rightarrow R$.

Exercise 8.2

Create an ImageJ plugin that shows the color table of an 8-bit indexed image as a new image with 16×16 rectangular color fields. Mark all unused color table entries in a suitable way. Look at Prog. 8.3 as a starting point.

¹³ It may seem a paradox that, although the RGB histogram is usually much larger than the image itself, the histogram is not sufficient in general to reconstruct the original image.

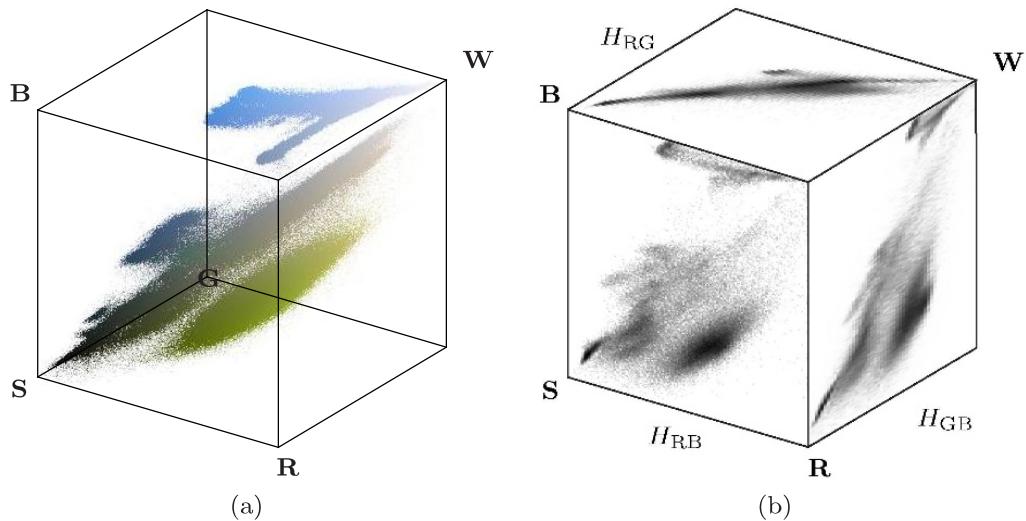


Figure 8.22 Two-dimensional RGB histogram projections. Three-dimensional RGB cube illustrating an image’s color distribution (a). The color points indicate the corresponding pixel colors and not the color frequency. The combined histograms for red-green (H_{RG}), red-blue (H_{RB}), and green-blue (H_{GB}) are 2D projections of the 3D histogram. The corresponding image is shown in Fig. 8.9 (a).

```

1  static int[][] get2dHistogram
2          (ColorProcessor cp, int c1, int c2) {
3      // c1, c2: R = 0, G = 1, B = 2
4      int[] RGB = new int[3];
5      int[][] H = new int[256][256]; // histogram array H[c1][c2]
6
7      for (int v = 0; v < cp.getHeight(); v++) {
8          for (int u = 0; u < cp.getWidth(); u++) {
9              cp.getPixel(u, v, RGB);
10             int i = RGB[c1];
11             int j = RGB[c2];
12             // increment corresponding histogram cell
13             H[j][i]++;
14         }
15     }
16     return H;
17 }
```

Program 8.11 Method `get2dHistogram()` for computing a combined 2D color histogram. The color components (histogram axes) are specified by the parameters `c1` and `c2`. The color distribution `H` is returned as a two-dimensional `int` array. The method is defined in class `ColorStatistics` (Prog. 8.10).

Exercise 8.3

Show that a “desaturated” RGB pixel produced in the form $(r, g, b) \rightarrow (y, y, y)$, where y is the equivalent luminance value (see Eqn. (8.8)), has the luminance y as well.

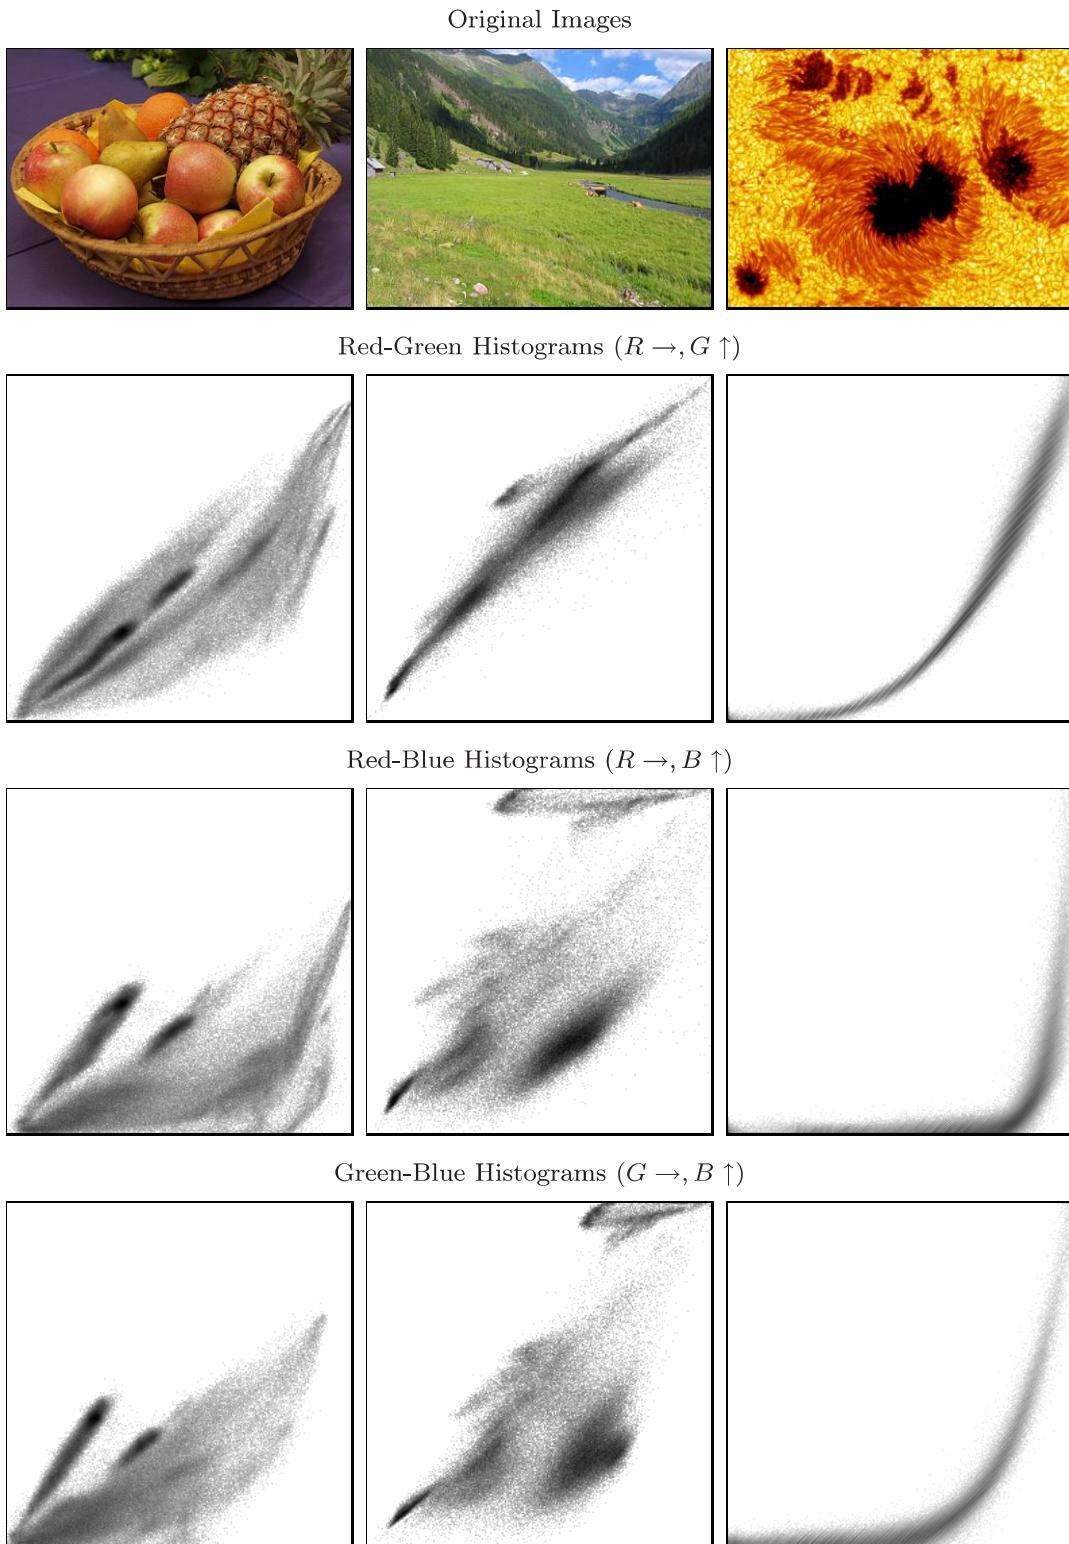


Figure 8.23 Combined color histogram examples. For better viewing, the images are inverted (dark regions indicate high frequencies) and the gray value corresponds to the logarithm of the histogram entries (scaled to the maximum entries).

Exercise 8.4

Extend the ImageJ plugin for desaturating color images in Prog. 8.5 such that the image is only modified inside the user-selected region of interest (ROI).

Exercise 8.5

Pseudocolors are sometimes used for displaying grayscale images (i. e., for viewing medical images with high dynamic range). Create an ImageJ plugin for converting 8-bit grayscale images to an indexed image with 256 colors, simulating the hues of glowing iron (from dark red to yellow and white).

Exercise 8.6

Determining the similarity between images of different sizes is a frequent problem (e. g., in the context of image data bases). Color statistics are commonly used for this purpose because they facilitate a coarse classification of images, such as landscape images, portraits, etc. However, two-dimensional color histograms (as described in Sec. 8.3.2) are usually too large and thus cumbersome to use for this purpose. A simple idea could be to split the 2D histograms or even the full RGB histogram into K regions (*bins*) and to combine the corresponding entries into a K -dimensional feature vector, which could be used for a coarse comparison. Develop a concept for such a procedure, and also discuss the possible problems.