

CHAPTER 12



Data Processing and Analysis

In the last several chapters we have covered the main topics of traditional scientific computing. These topics provide a foundation for most computational work. Starting with this chapter, we move on to explore data processing and analysis, statistics, and statistical modeling. As a first step in this direction, we look at the data analysis library `pandas`. This library provides convenient data structures for representing series and tables of data, and makes it easy to transform, split, merge, and convert data. These are important steps in the process¹ of cleansing raw data into a tidy form that is suitable for analysis. The `pandas` library builds on top of NumPy, and complements it with features that are particularly useful when handling data, such as labeled indexing, hierarchical indices, alignment of data for comparison and merging of datasets, handling of missing data, and much more. As such, the `pandas` library has become a de facto library for high-level data processing in Python, especially for statistics applications. The `pandas` library itself contains only limited support for statistical modeling (namely, linear regression). For more involved statistical analysis and modeling there are other packages available, such as `statmodels`, `patsy`, and `scikit-learn`, which we cover in later chapters. However, also for statistical modeling with these packages, `pandas` can still be used for data representation and preparation. The `pandas` library is therefore a key component in the software stack for data analysis with Python.

pandas The `pandas` library is a framework for data processing and analysis in Python. At the time of writing, the most recent version of `pandas` is 0.16.2. For more information about the `pandas` library, and its official documentation, see the project's web site at <http://pandas.pydata.org>.

The main focus of this chapter is to introduce basic features and usage of the `pandas` library. Toward the end of the chapter we also briefly explore the statistical visualization library `seaborn`, which is built on top of `Matplotlib`. This library provides quick and convenient graphing of data represented as `pandas` data structure (or NumPy arrays). Visualization is a very important part of exploratory data analysis, and the `pandas` library itself also provides functions for basic data visualization (which also builds on top of `Matplotlib`). The `seaborn` library takes this further, by providing additional statistical graphing capabilities and improved styling: The `seaborn` library is notable for generating good-looking graphics using defaults settings.

¹Also known as data munging or data wrangling.

■ **seaborn** The seaborn library is a visualization library for statistical graphics. It builds on Matplotlib and provides easy-to-use functions for common statistical graphs. At the time of writing, the most recent version of seaborn is 0.6.0. For more information about seaborn, and its official documentation, see the project's web site at: <http://stanford.edu/~mwaskom/software/seaborn>.

Importing Modules

In this chapter we mainly work with the `pandas` library, which we assume is imported under the name `pd`:

In [1]: `import pandas as pd`

We also require NumPy and Matplotlib, which we import as usual in the following way:

In [2]: `import numpy as np`

In [3]: `import matplotlib.pyplot as plt`

For more aesthetically pleasing default appearances of Matplotlib figures produced by the `pandas` library, we select an improved default style using the `pandas` function `set_option`:

In [4]: `pd.set_option('display.mpl_style', 'default')`

Later in this chapter we will also require to import the `seaborn` module, which we will import under the name `sns`, but for now we do not import this library since it alters the default appearance of graphics generated with Matplotlib.

Introduction to Pandas

The main focus of this chapter is the `pandas` library for data analysis, and we begin here with an introduction to this library. The `pandas` library mainly provides data structures and methods for representing and manipulating data. The two main data structures in `pandas` are the `Series` and `DataFrame` objects, which are used to represent data series and tabular data, respectively. Both of these objects have an index for accessing elements or rows in the data represented by the object. By default, the indices are integers starting from zero, like NumPy arrays, but it is also possible to use as index any sequence of identifiers.

Series

The merit of being able to index a data series with labels rather than integers is apparent even in the simplest of examples: Consider the following construction of a `Series` object. We give the constructor a list of integers, to create a `Series` object that represents the given data. Displaying the object in IPython reveals the data of the `Series` object together with the corresponding indices:

```
In [5]: s = pd.Series([909976, 8615246, 2872086, 2273305])
In [6]: s
Out[6]: 0      909976
         1      8615246
```

```
2      2872086
3      2273305
dtype: int64
```

The resulting object is a Series instance with data type (dtype) int64, and the elements are indexed by the integers 0, 1, 2, and 3. Using the index and values attributes, we can extract the underlying data for the index and the values stored in the series:

```
In [7]: s.index
Out[7]: Int64Index([0, 1, 2, 3], dtype='int64')
In [8]: s.values
Out[8]: array([ 909976, 8615246, 2872086, 2273305])
```

While using integer-indexed arrays or data series is a fully functional representation of the data, it is not descriptive. For example, if the data represents the population of four European capitals, it is convenient and descriptive to use the city names as indices rather than integers. With a Series object this is possible, and we can assign the index attribute of a Series object to a list with new indices to accomplish this. We can also set the name attribute of the Series object, to give it a descriptive name:

```
In [9]: s.index = ["Stockholm", "London", "Rome", "Paris"]
In [10]: s.name = "Population"
In [11]: s
Out[11]: Stockholm    909976
          London       8615246
          Rome         2872086
          Paris        2273305
Name: Population, dtype: int64
```

It is now immediately obvious what the data represents. Alternatively, we can also set the index and name attributes through keyword arguments to the Series object when it is created:

```
In [12]: s = pd.Series([909976, 8615246, 2872086, 2273305], name="Population",
...:                 index=["Stockholm", "London", "Rome", "Paris"])
```

While it is perfectly possible to store the data for the populations of these cities directly in a NumPy array, even in this simple example it is much clearer what the data represent when the data points are indexed with meaningful labels. The benefits of bringing the description of the data closer to the data are even greater when the complexity of the dataset increases.

We can access elements in a Series by indexing with the corresponding index (label), or directly through an attribute with the same name as the index (if the index label is a valid Python symbol name):

```
In [13]: s["London"]
Out[13]: 8615246
In [14]: s.Stockholm
Out[14]: 909976
```

Indexing a Series object with a list of indices gives a new Series object with a subset of the original data (corresponding to the provided list of indices):

```
In [15]: s[["Paris", "Rome"]]
Out[15]: Paris    2273305
          Rome    2872086
          Name: Population, dtype: int64
```

With a data series represented as a Series object, we can easily compute its descriptive statistics using the Series methods `count` (the number of data points), `median` (calculate the median), `mean` (calculate the mean value), `std` (calculate the standard deviation), `min` and `max` (minimum and maximum value), and the `quantile` (for calculating quantiles):

```
In [16]: s.median(), s.mean(), s.std()
Out[16]: (2572695.5, 3667653.25, 3399048.5005155364)
In [17]: s.min(), s.max()
Out[17]: (909976, 8615246)
In [18]: s.quantile(q=0.25), s.quantile(q=0.5), s.quantile(q=0.75)
Out[18]: (1932472.75, 2572695.5, 4307876.0)
```

All of the above are combined in the output of the `describe` method, which provides a summary of the data represented by a Series object:

```
In [19]: s.describe()
Out[19]: count      4.000000
          mean     3667653.250000
          std      3399048.500516
          min      909976.000000
          25%     1932472.750000
          50%     2572695.500000
          75%     4307876.000000
          max     8615246.000000
          Name: Population, dtype: float64
```

Using the `plot` method, we can quickly and easily produce graphs that visualize the data in a Series object. The pandas library uses Matplotlib to produce graphs, and we can optionally pass a Matplotlib Axes instance to the `plot` method via the `ax` argument. The type of the graph is specified using the `kind` argument (valid options are `line`, `hist`, `bar`, `barh`, `box`, `kde`, `density`, `area` and `pie`). See Figure 12-1.

```
In [20]: fig, axes = plt.subplots(1, 4, figsize=(12, 3))
...: s.plot(ax=axes[0], kind='line', title='line')
...: s.plot(ax=axes[1], kind='bar', title='bar')
...: s.plot(ax=axes[2], kind='box', title='box')
...: s.plot(ax=axes[3], kind='pie', title='pie')
```

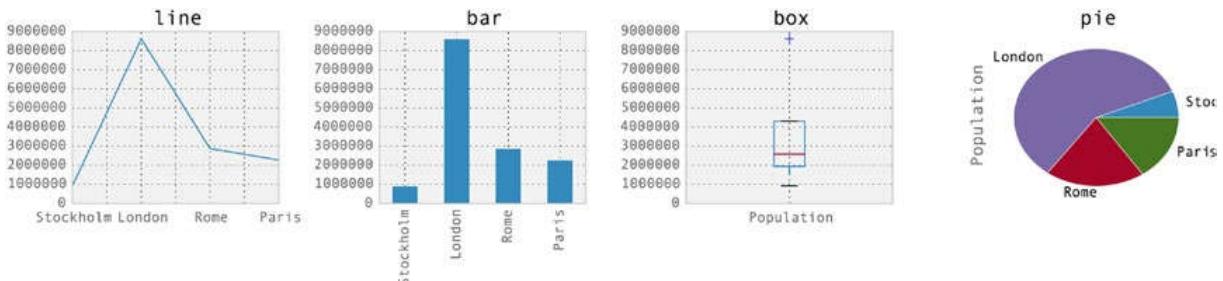


Figure 12-1. Examples of plot styles that can be produced with pandas using the `Series.plot` method

DataFrame

As we have seen in the previous examples, a pandas `Series` object provides a convenient container for one-dimensional arrays, which can use descriptive labels for the elements, and which provides quick access to descriptive statistics and visualization. For higher-dimensional arrays (mainly two-dimensional arrays or tables), the corresponding data structure is the pandas `DataFrame` object. It can be viewed as a collection of `Series` objects with a common index.

There are numerous ways to initialize a `DataFrame`. For simple examples, the easiest way is to pass a nested Python list or dictionary to the constructor of the `DataFrame` object. For example, consider an extension of the dataset we used in the previous section, where in addition to the population of each city we also include a column that specifies which state each city belongs to. We can create the corresponding `DataFrame` object in the following way:

```
In [21]: df = pd.DataFrame([[909976, "Sweden"],
...:                      [8615246, "United Kingdom"],
...:                      [2872086, "Italy"],
...:                      [2273305, "France"]])
In [22]: df
Out[22]:
```

	0	1
0	909976	Sweden
1	8615246	United Kingdom
2	2872086	Italy
3	2273305	France

The result is tabular data structure with rows and columns. Like with a `Series` object, we can use labeled indexing for rows by assigning a sequence of labels to the `index` attribute, and, in addition, we can set the `columns` attribute to a sequence of labels for the columns:

```
In [23]: df.index = ["Stockholm", "London", "Rome", "Paris"]
In [24]: df.columns = ["Population", "State"]
```

In [25]: df
Out[25]:

	Population	State
Stockholm	909976	Sweden
London	8615246	United Kingdom
Rome	2872086	Italy
Paris	2273305	France

The index and columns attributes can also be set using the corresponding keyword arguments to the DataFrame object when it is created:

```
In [26]: df = pd.DataFrame([[909976, "Sweden"],
...:                      [8615246, "United Kingdom"],
...:                      [2872086, "Italy"],
...:                      [2273305, "France"]],
...:                      index=["Stockholm", "London", "Rome", "Paris"],
...:                      columns=["Population", "State"])
```

An alternative way to create the same data frame, which sometimes can be more convenient, is to pass a dictionary with column titles as keys and column data as values:

```
In [27]: df = pd.DataFrame({"Population": [909976, 8615246, 2872086, 2273305],
...:                         "State": ["Sweden", "United Kingdom", "Italy", "France"]},
...:                         index=["Stockholm", "London", "Rome", "Paris"])
```

As before, the underlying data in a DataFrame can be obtained as a NumPy array using the values attribute, and the index and column arrays through the index and columns attributes, respectively. Each column in a data frame can be accessed using the column name as attribute (or, alternatively, by indexing with the column label, for example df["Population"]):

```
In [28]: df.Population
Out[28]: Stockholm    909976
          London      8615246
          Rome        2872086
          Paris       2273305
          Name: Population, dtype: int64
```

The result of extracting a column from a DataFrame is a new Series object, which we can process and manipulate with the methods discussed in the previous section. Rows of a DataFrame instance can be accessed using the ix indexer attribute. Indexing this attribute also results in a Series object, which corresponds to a row of the original data frame:

```
In [29]: df.ix["Stockholm"]
Out[29]: Population    909976
          State        Sweden
          Name: Stockholm, dtype: object
```

Passing a list of row labels to the `ix` indexer results in a new `DataFrame` that is a subset of the original `DataFrame`, containing only the selected rows:

```
In [30]: df.ix[["Paris", "Rome"]]
Out[30]:
```

	Population	State
Paris	2273305	France
Rome	2872086	Italy

The `ix` indexer can also be used to select both rows and columns simultaneously, by passing first a row label (or list thereof), and second a column label (or list thereof). The result is a `DataFrame`, a `Series`, or an element value, depending on the number of columns and rows that are selected:

```
In [31]: df.ix[["Paris", "Rome"], "Population"]
Out[31]: Paris    2273305
          Rome    2872086
          Name: Population, dtype: int64
```

We can compute descriptive statistics using the same methods as we already used for `Series` objects. When invoking those methods (`mean`, `std`, `median`, `min`, `max`, etc.) for a `DataFrame`, the calculation is performed for each column with numerical data types:

```
In [32]: df.mean()
Out[32]: Population    3667653.25
          dtype: float64
```

In this case, only one of the two columns has a numerical data type (the one named `Population`). Using the `DataFrame` method `info` and the attribute `dtype`, we can obtain a summary of the content in a `DataFrame`, and the data types of each column:

```
In [33]: df.info()
<class 'pandas.core.frame.DataFrame'>
Index: 4 entries, Stockholm to Paris
Data columns (total 2 columns):
Population    4 non-null int64
State         4 non-null object
dtypes: int64(1), object(1)
memory usage: 96.0+ bytes
In [34]: df.dtypes
Out[34]: Population    int64
          State        object
          dtype: object
```

The real advantages of using `pandas` emerge when dealing with larger and more complex datasets than the examples we have used so far. Such data can rarely be defined as explicit lists or dictionaries, which can be passed to the `DataFrame` initializer. A more common situation is that the data must be read from a file, or some other external source. The `pandas` library supports numerous methods for reading data from files of different formats. Here we use the `read_csv` function to read in data and create a `DataFrame` object

from a CSV file.² This function accepts a large number of optional arguments for tuning its behavior. See the docstring `help(pd.read_csv)` for details. Some of the most useful arguments are `header` (specifies which row, if any, contains a header with column names), `skiprows` (number of rows to skip before starting to read data, or a list of line numbers of lines to skip), `delimiter` (the character that is used as a delimiter between column values), `encoding` (the name of the encoding used in the file, for example `utf-8`), and `nrows` (number of rows to read). The first and only mandatory argument to the `pd.read_csv` function is a filename or an URL to the data source. For example, to read in a dataset stored in a file called `european_cities.csv`³, of which the first five lines are shown below, we can simply call `pd.read_csv("european_cities.csv")`, since the default delimiter is `,` and the header is by default taken from the first line. However, we could also write out all these options explicitly:

```
In [35]: !head -n 5 european_cities.csv
Rank,City,State,Population,Date of census
1,London, United Kingdom,"8,615,246",1 June 2014
2,Berlin, Germany,"3,437,916",31 May 2014
3,Madrid, Spain,"3,165,235",1 January 2014
4,Rome, Italy,"2,872,086",30 September 2014
In [36]: df_pop = pd.read_csv("european_cities.csv",
...:                         delimiter=",", encoding="utf-8", header=0)
```

This dataset is similar to the example data we used earlier in this chapter, but here there are additional columns and many more rows for other cities. Once a dataset is read into a `DataFrame` object, it is useful to start by inspecting the summary given by the `info` method, to begin forming an idea of the properties of the dataset.

```
In [37]: df_pop.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 105 entries, 0 to 104
Data columns (total 5 columns):
Rank           105 non-null int64
City           105 non-null object
State          105 non-null object
Population     105 non-null object
Date of census 105 non-null object
dtypes: int64(1), object(4) memory usage: 4.9+ KB
```

Here we see that there are 105 rows in this dataset, and that it has five columns. Only the `Rank` column is of numerical data type. In particular, the `Population` column is not yet of numeric data type because its values are of the format `"8,615,246,"` and is therefore interpreted as string values by the `read_csv` function. It is also informative to display a tabular view of the data. However, this dataset is too large to display in full, and in situations like this the `head` and `tail` methods are handy for creating a truncated dataset containing the first few and last few rows, respectively. Both of these functions take an optional argument that specifies how many rows to include in the truncated `DataFrame`. Note also that `df.head(n)` is equivalent to `df[:n]`, where `n` is an integer.

²CSV, or comma-separated values, is a common text format where rows are stored in lines and columns are separated by a comma (or some other text delimiter). See Chapter 18 for more details about this and other file formats.

³This dataset was obtained from the Wiki page: http://en.wikipedia.org/wiki/Largest_cities_of_the_European_Union_by_population_within_city_limits.

In [38]: df_pop.head()

Out[38]:

	Rank	City	State	Population	Date of census
0	1	London	United Kingdom	8,615,246	1 June 2014
1	2	Berlin	Germany	3,437,916	31 May 2014
2	3	Madrid	Spain	3,165,235	1 January 2014
3	4	Rome	Italy	2,872,086	30 September 2014
4	5	Paris	France	2,273,305	1 January 2013

Displaying a truncated DataFrame gives a good idea of how the data looks, and what remains to be done before the data is ready for analysis. It is common to have to transform columns in one way or another, and to reorder the table by sorting by a specific column, or by ordering the index. In the following we explore some methods for modifying DataFrame objects. First of all, we can create new columns and update columns in a DataFrame simply by assigning a Series object to the DataFrame indexed by the column name, and we can delete columns using the Python `del` keyword.

The `apply` method is a powerful tool to transform the content in a column. It creates and returns a new Series object for which a function passed to `apply` has been applied to each element in the original column. For example, we can use the `apply` method to transform the elements in the `Population` column from strings to integers, by passing a lambda function that removes the `,` characters from the strings and casts the results to an integer. Here we assign the transformed column to a new column with name `NumericPopulation`. Using the same method, we also tidy up the `State` values by removing extra white spaces in its elements using the string method `strip`.

```
In [39]: df_pop["NumericPopulation"] = df_pop.Population.apply(
...:     lambda x: int(x.replace(",","")))
In [40]: df_pop["State"].values[:3] # contains extra white spaces
Out[40]: array([' United Kingdom', ' Germany', ' Spain'], dtype=object)
In [41]: df_pop["State"] = df_pop["State"].apply(lambda x: x.strip())
In [42]: df_pop.head()
Out[42]:
```

	Rank	City	State	Population	Date of census	NumericPopulation
0	1	London	United Kingdom	8,615,246	1 June 2014	8615246
1	2	Berlin	Germany	3,437,916	31 May 2014	3437916
2	3	Madrid	Spain	3,165,235	1 January 2014	3165235
3	4	Rome	Italy	2,872,086	30 September 2014	2872086
4	5	Paris	France	2,273,305	1 January 2013	2273305

Inspecting the data types of the columns in the updated DataFrame confirms that the new column `NumericPopulation` is indeed of integer type (while the `Population` column is unchanged):

```
In [43]: df_pop.dtypes
Out[43]: Rank          int64
          City         object
          State        object
```

```
Population      object
Date of census    object
NumericPopulation int64
dtype: object
```

We may also need to change the index to one of the columns of the DataFrame. In the current example, we may want to use the City column as index. We can accomplish this using the `set_index` method, which takes as argument the name of the column to use as index. The result is a new DataFrame object, and the original DataFrame is unchanged. Furthermore, using the `sort_index` method we can sort the data frame with respect to the index:

```
In [44]: df_pop2 = df_pop.set_index("City")
In [45]: df_pop2 = df_pop2.sort_index()
In [46]: df_pop2.head()
Out[46]:
```

	Rank	State	Population	Date of census	NumericPopulation
City					
Aarhus	92	Denmark	326,676	1 October 2014	326676
Alicante	86	Spain	334,678	1 January 2012	334678
Amsterdam	23	Netherlands	813,562	31 May 2014	813562
Antwerp	59	Belgium	510,610	1 January 2014	510610
Athens	34	Greece	664,046	24 May 2011	664046

The `sort_index` method also accepts a list of column names, in which case a hierarchical index is created. A hierarchical index uses tuples of index labels to address rows in the data frame. We can use the `sortlevel` method, which takes an integer `n` as argument, to sort the rows in a DataFrame according to the `n`th level of the hierarchical index. In the following example we create a hierarchical index with State and City as indices, and we use the `sortlevel` method to sort by the first index (State):

```
In [47]: df_pop3 = df_pop.set_index(["State", "City"]).sortlevel(0)
In [48]: df_pop3.head(7)
Out[48]:
```

		Rank	Population	Date of census
State	City			
Austria	Vienna	7	1794770	1 January 2015
Belgium	Antwerp	59	510610	1 January 2014
	Brussels	16	1175831	1 January 2014
Bulgaria	Plovdiv	84	341041	31 December 2013
	Sofia	14	1291895	14 December 2014
	Varna	85	335819	31 December 2013
Croatia	Zagreb	24	790017	31 March 2011

A DataFrame with a hierarchical index can be partially indexed using only its zeroth-level index (`df3.ix["Sweden"]`), or completely indexed using a tuple of all hierarchical indices (`df3.ix[("Sweden", "Gothenburg")]`):

In [49]: `df_pop3.ix["Sweden"]`

Out[49]:

	Rank	Population	Date of census	NumericPopulation
City				
Gothenburg	53	528,014	31 March 2013	528014
Malmö	102	309,105	31 March 2013	309105
Stockholm	20	909,976	31 January 2014	909976

In [50]: `df_pop3.ix[("Sweden", "Gothenburg")]`

Out[50]: Rank 53

Population 528,014

Date of census 31 March 2013

NumericPopulation 528014

Name: (Sweden, Gothenburg), dtype: object

If we want to sort by a column rather than the index, we can use the `sort` method. It takes a column name, or a list of column names, with respect to which the DataFrame is to be sorted. It also accepts the keyword argument `ascending`, which is a Boolean or a list of Boolean values that specifies whether the corresponding column is to be sorted in ascending or descending order:

In [51]: `df_pop.set_index("City").sort(["State", "NumericPopulation"], ascending=[False, True]).head()`

Out[51]:

	Rank	State	Population	Date of census	NumericPopulation
City					
Nottingham	103	United Kingdom	308,735	30 June 2012	308735
Wirral	97	United Kingdom	320,229	30 June 2012	320229
Coventry	94	United Kingdom	323,132	30 June 2012	323132
Wakefield	91	United Kingdom	327,627	30 June 2012	327627
Leicester	87	United Kingdom	331,606	30 June 2012	331606

With categorical data such as the `State` column, it is frequently of interest to summarize how many values of each category a column contains. Such counts can be computed using the `value_counts` method (of the Series object). For example, to count the number of cities each country has in the list of the 105 largest cities in Europe, we can use:

In [52]: `city_counts = df_pop.State.value_counts()`

In [53]: `city_counts.head()`

Out[53]: Germany 19
United Kingdom 16

```
Spain           13
Poland          10
Italy            10
dtype: int64
```

In this example, we see from the results that the state with the largest number of cities in the list is Germany, with 19 cities, followed by the United Kingdom with 16 cities, and so on. A related question is how large the total population of all cities within a state. To answer this type of question we can proceed in two ways: first, we can create a hierarchical index using State and City, and use the sum method to reduce the DataFrame along the one of the indices. In this case, we want to sum over all entries within the index level State, so we can use sum(level="State"), which eliminates the City index. For presentation we also sort the resulting DataFrame in descending order of the column NumericPopulation:

```
In [54]: df_pop3 = df_pop[["State", "City", "NumericPopulation"]].set_index(["State", "City"])
In [55]: df_pop4 = df_pop3.sum(level="State").sort("NumericPopulation", ascending=False)
In [56]: df_pop4.head()
Out[56]:
```

	NumericPopulation
State	
United Kingdom	16011877
Germany	15119548
Spain	10041639
Italy	8764067
Poland	6267409

Second, we can obtain the same results using the groupby method, which allows us to group rows of a DataFrame by the values of a given column, and apply a reduction function on the resulting object (for example, sum, mean, min, max, etc.). The result is a new DataFrame with the grouped-by column as index. Using this method we can compute the total population in the 105 cities, grouped by state, in the following way.

```
In [57]: df_pop5 = (df_pop.drop("Rank", axis=1)
...:                 .groupby("State").sum()
...:                 .sort("NumericPopulation", ascending=False))
```

Note that here we also used the drop method to remove the Rank column (hence the axis=1, use axis=0 to drop rows) from the DataFrame (since it is not meaningful to aggregate the rank by summation). Finally, we use the plot method of the Series object to plot bar graphs for the city count and the total population. The results are shown in Figure 12-2.

```
In [58]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
...: city_counts.plot(kind='barh', ax=ax1)
...: ax1.set_xlabel("# cities in top 105")
...: df_pop5.NumericPopulation.plot(kind='barh', ax=ax2)
...: ax2.set_xlabel("Total pop. in top 105 cities")
```

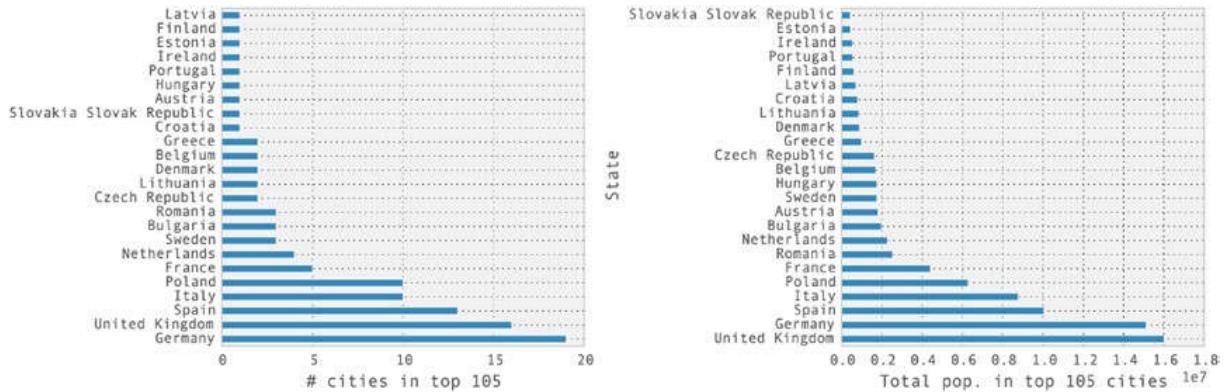


Figure 12-2. The number of cities in the list of the top 105 most populated cities in Europe (left) and the total population in those cities (right), grouped by state

Time Series

Time series are a common form of data in which a quantity is given, for example, at regularly or irregularly spaced timestamps, or for fixed or variable time spans (periods). In pandas, there are dedicated data structures for representing these types of data. Series and DataFrame can have both columns and indices with data types describing timestamps and time spans. When dealing with temporal data it is particularly useful to be able to index the data with time data types. Using pandas time-series indexers, DatetimeIndex and PeriodIndex, we can carry out many common date, time, period, and calendar operations, such as selecting time ranges, shifting and resampling of the data points in a time series.

To generate a sequence of dates that can be used as an index in a pandas Series or DataFrame objects, we can, for example, use the `date_range` function. It takes the starting point as a date and time string (or, alternatively, a `datetime` object from the Python standard library) as a first argument, and the number of elements in the range can be set using the `periods` keyword argument:

```
In [59]: pd.date_range("2015-1-1", periods=31)
Out[59]: <class 'pandas.tseries.index.DatetimeIndex'>
          [2015-01-01, ..., 2015-01-31]
Length: 31, Freq: D, Timezone: None
```

To specify the frequency of the timestamps (which defaults to one day), we can use the `freq` keyword argument, and instead of using `periods` to specify the number of points, we can give both starting and ending points as date and time strings (or `datetime` objects) as first and second argument. For example, to generate hourly timestamps between 00:00 and 12:00 on 2015-01-01, we can use:

```
In [60]: pd.date_range("2015-1-1 00:00", "2015-1-1 12:00", freq="H")
Out[60]: <class 'pandas.tseries.index.DatetimeIndex'>
          [2015-01-01 00:00:00, ..., 2015-01-01 12:00:00]
Length: 13, Freq: H, Timezone: None
```

The `date_range` function returns an instance of `DatetimeIndex`, which can be used, for example, as an index for a `Series` or `DataFrame` object:

```
In [61]: ts1 = pd.Series(np.arange(31), index=pd.date_range("2015-1-1", periods=31))
In [62]: ts1.head()
Out[62]: 2015-01-01    0
          2015-01-02    1
          2015-01-03    2
          2015-01-04    3
          2015-01-05    4
Freq: D, dtype: int64
```

The elements of a `DatetimeIndex` object can, for example, be accessed using indexing with date and time strings. An element in a `DatetimeIndex` is of the type `Timestamp`, which is a pandas object that extends the standard Python `datetime` object (see the `datetime` module in the Python standard library).

```
In [63]: ts1["2015-1-3"]
Out[63]: 2
In [64]: ts1.index[2]
Out[64]: Timestamp('2015-01-03 00:00:00', offset='D')
```

In many aspects, a `Timestamp` and `datetime` object are interchangeable, and the `Timestamp` class have, like the `datetime` class, attributes for accessing time fields such as `year`, `month`, `day`, `hour`, `minute`, and so on. However, a notable difference between `Timestamp` and `datetime` is that `Timestamp` store a timestamp with nanoseconds resolution, while a `datetime` object only uses microsecond resolution.

```
In [65]: ts1.index[2].year, ts1.index[2].month, ts1.index[2].day
Out[65]: (2015, 1, 3)
In [66]: ts1.index[2].nanosecond
Out[66]: 0
```

We can convert a `Timestamp` object to a standard Python `datetime` object using the `to_pydatetime` method:

```
In [67]: ts1.index[2].to_pydatetime()
Out[67]: datetime.datetime(2015, 1, 3, 0, 0)
```

and we can use a list of `datetime` objects to create a pandas time series:

```
In [68]: import datetime
In [69]: ts2 = pd.Series(np.random.rand(2),
...:                     index=[datetime.datetime(2015, 1, 1), datetime.datetime(2015, 2, 1)])
In [70]: ts2
Out[70]: 2015-01-01    0.683801
          2015-02-01    0.916209
          dtype: float64
```

Data that is defined for sequences of time spans can be represented using Series and DataFrame objects that are indexed using the PeriodIndex class. We can construct an instance of the PeriodIndex class explicitly by passing a list of Period objects, and then specify it as an index when creating a Series or DataFrame object:

```
In [71]: periods = pd.PeriodIndex([pd.Period('2015-01'),
...:                               pd.Period('2015-02'),
...:                               pd.Period('2015-03')])
In [72]: ts3 = pd.Series(np.random.rand(3), index=periods)
In [73]: ts3
Out[73]: 2015-01    0.969817
          2015-02    0.086097
          2015-03    0.016567
          Freq: M, dtype: float64
In [74]: ts3.index
Out[74]: <class 'pandas.tseries.period.PeriodIndex'>
          [2015-01, ..., 2015-03]
          Length: 3, Freq: M
```

We can also convert a Series or DataFrame object indexed by a DatetimeIndex object to a PeriodIndex using the `to_period` method (which takes an argument that specifies the period frequency, here 'M' for month):

```
In [75]: ts2.to_period('M')
Out[75]: 2015-01    0.683801
          2015-02    0.916209
          Freq: M, dtype: float64
```

In the remaining part of this section we explore select features of pandas time series through examples. We look at the manipulation of two time series that contain sequences of temperature measurements at given timestamps. We have one dataset for an indoor temperature sensor, and one dataset for an outdoors temperature sensor, both with observations approximately every 10 minutes during most of 2014. The two data files, `temperature_indoor_2014.tsv` and `temperature_outdoor_2014.tsv`, are TSV (tab-separated values, a variant of the CSV format) files with two columns: the first column contains UNIX timestamps (seconds since Jan 1, 1970), and the second column is the measured temperature in degree Celsius. For example, the first five lines in the outdoor dataset are:

```
In [76]: !head -n 5 temperature_outdoor_2014.tsv
1388530986    4.380000
1388531586    4.250000
1388532187    4.190000
1388532787    4.060000
1388533388    4.060000
```

We can read the data files using `read_csv` by specifying that the delimiter between columns is the TAB character: `delimiter="\t"`. When reading the two files we also explicitly specify the column names using the `names` keyword argument, since the files in this example do not have header lines with the column names.

```
In [77]: df1 = pd.read_csv('temperature_outdoor_2014.tsv', delimiter='\t',
...:                         names=["time", "outdoor"])
In [78]: df2 = pd.read_csv('temperature_indoor_2014.tsv', delimiter='\t',
...:                         names=["time", "indoor"])
```

Once we have created DataFrame objects for the time series data, it is informative to inspect the data by displaying the first few lines:

```
In [79]: df1.head()
Out[79]:
```

	time	outdoor
0	1388530986	4.38
1	1388531586	4.25
2	1388532187	4.19
3	1388532787	4.06
4	1388533388	4.06

The next step toward a meaningful representation of the time series data is to convert the UNIX timestamps to date and time objects using `to_datetime` with the `unit="s"` argument. Furthermore, we localize the timestamps (assigning a time zone) using `tz_localize` and convert the time zone attribute to the Europe/Stockholm time zone using `tz_convert`. We also set the `time` column as index using `set_index`:

```
In [80]: df1.time = (pd.to_datetime(df1.time.values, unit="s")
...:                 .tz_localize('UTC').tz_convert('Europe/Stockholm'))
In [81]: df1 = df1.set_index("time")
In [82]: df2.time = (pd.to_datetime(df2.time.values, unit="s")
...:                 .tz_localize('UTC').tz_convert('Europe/Stockholm'))
In [83]: df2 = df2.set_index("time")
In [84]: df1.head()
Out[84]:
```

	outdoor
time	
2014-01-01 00:03:06+01:00	4.38
2014-01-01 00:13:06+01:00	4.25
2014-01-01 00:23:07+01:00	4.19
2014-01-01 00:33:07+01:00	4.06
2014-01-01 00:43:08+01:00	4.06

Displaying the first few rows of the data frame for the outdoor temperature dataset shows that the index now indeed is a date and time object. As we will see examples of in the following, having the index of a time series represented as proper date and time objects (in contrast to using for example integers representing the UNIX timestamps) allows us to easily perform many time-oriented operations. Before we proceed to explore the data in more detail, we first plot the two time series to obtain an idea of how the data looks like.

For this we can use the `DataFrame.plot` method, and the results are shown in Figure 12-3. Note that there is data missing for a part of August. Imperfect data is a common problem, and handling missing data in a suitable manner is an important part of the mission statement of the pandas library.

```
In [85]: fig, ax = plt.subplots(1, 1, figsize=(12, 4))
...: df1.plot(ax=ax)
...: df2.plot(ax=ax)
```

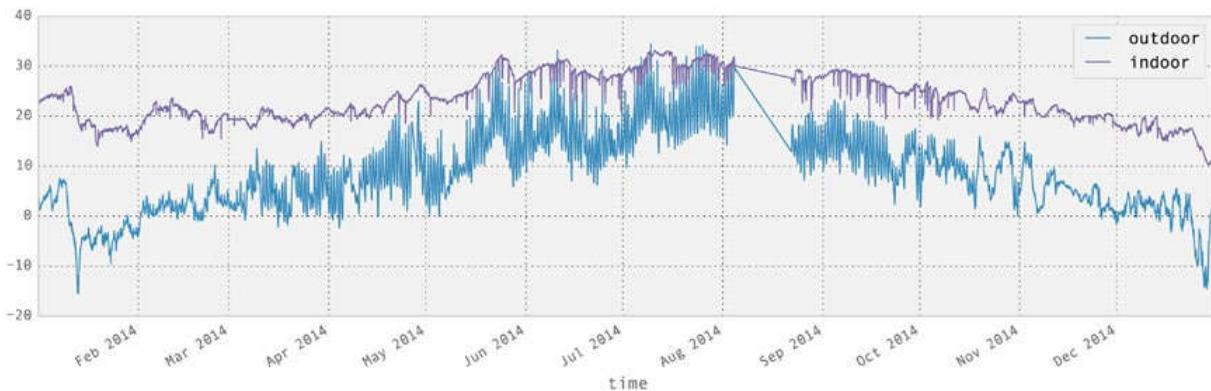


Figure 12-3. Plot of the time series for indoors and outdoors temperatures

It is also illuminating to display the result of the `info` method of the `DataFrame` object. Doing so tells us that there are nearly 50000 data points in this data set, and that it contains data points starting at 2014-01-01 00:03:06 and ending at 2014-12-30 23:56:35:

```
In [86]: df1.info()
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 49548 entries, 2014-01-01 00:03:06+01:00 to 2014-12-30 23:56:35+01:00
Data columns (total 1 columns):
outdoor    49548 non-null float64
dtypes: float64(1) memory usage: 774.2 KB
```

A common operation on time series is to select and extract parts of the data. For example, from the full dataset that contains data for all of 2014, we may be interested in selecting out and analyze only the data for the month of January. In pandas, we can accomplish this in a number of ways. For example, we can use Boolean indexing of a `DataFrame` to create a `DataFrame` for a subset of the data. To create the Boolean indexing mask that selects the data for January, we can take advantage of the pandas time series features that allows us to compare the time series index with string representations of a date and time. In the following code, the expressions like `df1.index >= "2014-1-1"`, where `df1.index` is a time `DateTimeIndex` instance, results in a Boolean NumPy array that can be used as a mask to select the desired elements.

```
In [87]: mask_jan = (df1.index >= "2014-1-1") & (df1.index < "2014-2-1")
In [88]: df1_jan = df1[mask_jan]
In [89]: df1_jan.info()
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 4452 entries, 2014-01-01 00:03:06+01:00 to 2014-01-31 23:56:58+01:00
Data columns (total 1 columns):
outdoor    4452 non-null float64
dtypes: float64(1) memory usage: 69.6 KB
```

Alternatively, we can use slice syntax directly with date and time strings:

```
In [90]: df1_jan = df1["2014-1-1":"2014-1-31"]
```

The results are two DataFrame objects, `df1_jan` and `df2_jan`, that contains data only for the month of January. Plotting this subset of the original data using the `plot` method results in the graph shown in Figure 12-4.

```
In [91]: fig, ax = plt.subplots(1, 1, figsize=(12, 4))
....: df1_jan.plot(ax=ax)
....: df2_jan.plot(ax=ax)
```

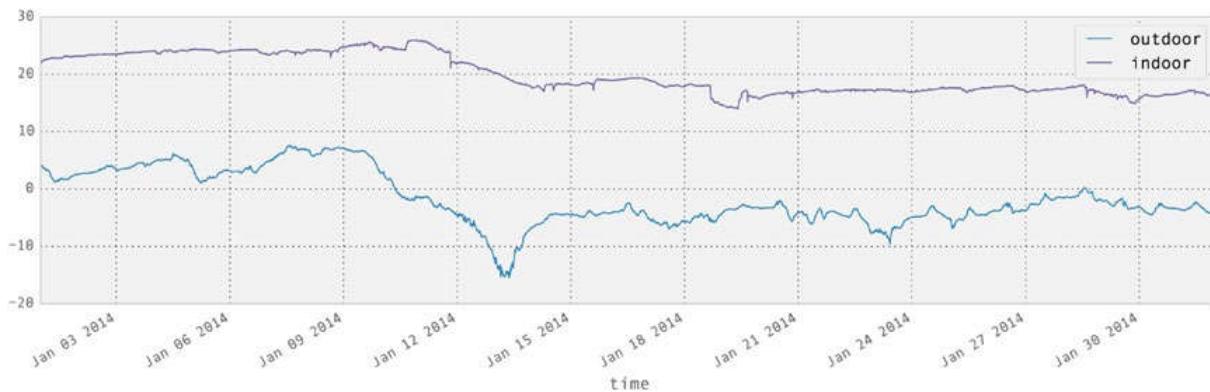


Figure 12-4. Plot of the time series for indoors and outdoors temperatures for a selected month (January)

Like the `datetime` class in Python's standard library, the `Timestamp` class that is used in pandas to represent time values has attributes for accessing fields such as `year`, `month`, `day`, `hour`, `minute`, and so on. These fields are particularly useful when processing time series. For example, if we wish to calculate the average temperature for each month of the year, we can begin by creating a new column `month`, which we assign to the `month` field of the `Timestamp` values of the `DatetimeIndex` indexer. To extract the `month` field from each `Timestamp` value, we first call `reset_index` to convert the index to a column in the data frame (in which case the new `DataFrame` object falls back to using an integer index), after which we can use the `apply` function on the newly created `time` column.⁴

```
In [92]: df1_month = df1.reset_index()
In [93]: df1_month["month"] = df1_month.time.apply(lambda x: x.month)
In [94]: df1_month.head()
Out[94]:
```

	time	outdoor	month
0	2014-01-01 00:03:06+01:00	4.38	1
1	2014-01-01 00:13:06+01:00	4.25	1
2	2014-01-01 00:23:07+01:00	4.19	1
3	2014-01-01 00:33:07+01:00	4.06	1
4	2014-01-01 00:43:08+01:00	4.06	1

⁴We can also directly use the `month` method of the `DatetimeIndex` index object, but for the sake of demonstration we use a more explicit approach here.

Next, we can group the DataFrame by the new month field, and aggregate the grouped values using the mean function for computing the average within each group.

```
In [95]: df1_month = df1_month.groupby("month").aggregate(np.mean)
In [96]: df2_month = df2.reset_index()
In [97]: df2_month["month"] = df2_month.time.apply(lambda x: x.month)
In [98]: df2_month = df2_month.groupby("month").aggregate(np.mean)
```

After having repeated the same process for the second DataFrame (indoor temperatures), we can combine df1_month and df2_month into a single DataFrame using the join method:

```
In [99]: df_month = df1_month.join(df2_month)
In [100]: df_month.head(3)
Out[100]:
```

	outdoor	indoor
time		
1	-1.776646	19.862590
2	2.231613	20.231507
3	4.615437	19.597748

In only a few lines of code, we have here leveraged some of the many data processing capabilities of pandas to transform and compute with the data. It is often the case that there are many different ways to combine the tools provided by pandas to do the same, or a similar, analysis. For the current example, we can do the whole process in a single line of code, using the to_period and groupby methods, and the concat function (which like join combines DataFrame into a single DataFrame):

```
In [101]: df_month = pd.concat([df.to_period("M").groupby(level=0).mean() for df in [df1, df2]],
...:                           axis=1)
In [102]: df_month.head(3)
Out[102]:
```

	outdoor	indoor
time		
2014-01	-1.776646	19.862590
2014-02	2.231613	20.231507
2014-03	4.615437	19.597748

To visualize the results, we plot the average monthly temperatures as a bar plot and a box plot using the DataFrame method plot. The result is shown in Figure 12-5.

```
In [103]: fig, axes = plt.subplots(1, 2, figsize=(12, 4))
...: df_month.plot(kind='bar', ax=axes[0])
...: df_month.plot(kind='box', ax=axes[1])
```

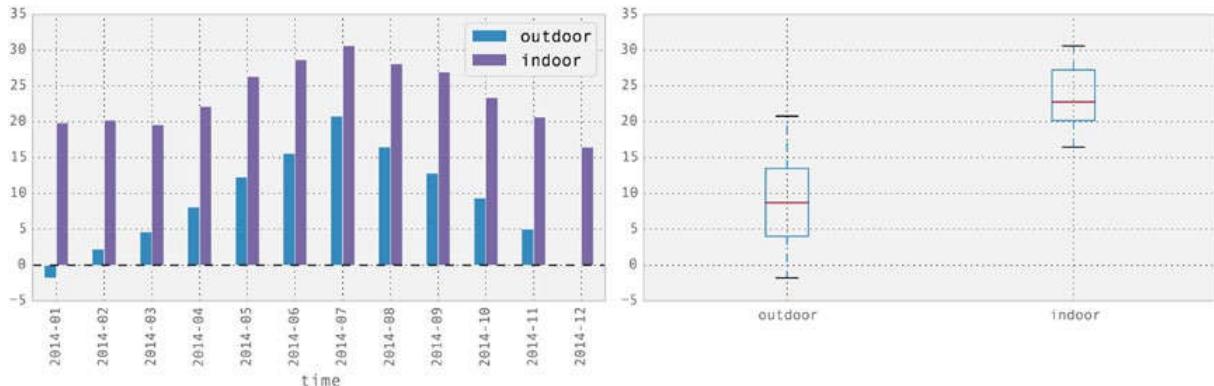


Figure 12-5. Average indoor and outdoor temperatures per month (left), and a boxplot for monthly indoor and outdoors temperature (right)

Finally, a very useful feature of the pandas time series objects is the ability to up- and down-sample the time series using the `resample` method. Resampling means that the number of data points in a time series is changed. It can be either increased (upsampling) or decreased (downsampling). For upsampling, we need to choose a method for filling in the missing values, and for downsampling we need to choose a method for aggregating multiple sample points between each new sample point. The `resample` method expects as first argument a string that specifies the new period of data points in the resampled time series. For example, the string `H` represents a period of one hour, the string `D` one day, the string `M` one month, and so on.⁵ We can also combine these in simple expressions, such as `7D`, which denotes a time period of seven days. Optionally, we can also use the `how` argument to specify how to aggregate values in the case of downsampling, and the `fill_method` argument to specify a method for filling in the value of new data points in the case of upsampling.

To illustrate the use for the `resample` method, consider the previous two time series with temperature data. The original sampling frequency is roughly 10 minutes, which amounts to a lot of data points over the period of a year. For plotting purposes, or if we want to compare the 2 time series, which are sampled at slightly different timestamps, it is often necessary to down-sample the original data. This can give less busy graphs, and regularly spaced time series that readily can be compared to each other. In the following code we resample the outdoors temperature time series to four different sampling frequencies, and plot the resulting time series. We also resample both the outdoor and indoor time series to daily averages that we subtract to obtain the daily average temperature difference between indoors and outdoors throughout the year. These types of manipulations are very handy when dealing time series, and it is one of the many areas in which the pandas library really shines. See Figure 12-6.

```
In [104]: df1_hour = df1.resample("H")
In [105]: df1_hour.columns = ["outdoor (hourly avg.)"]
In [106]: df1_day = df1.resample("D")
In [107]: df1_day.columns = ["outdoor (daily avg.)"]
In [108]: df1_week = df1.resample("7D")
In [109]: df1_week.columns = ["outdoor (weekly avg.)"]
In [110]: df1_month = df1.resample("M")
In [111]: df1_month.columns = ["outdoor (monthly avg.)"]
In [112]: df_diff = (df1.resample("D").outdoor - df2.resample("D").indoor)
```

⁵There are a large number of available time-unit codes. See the sections on “Offset aliases” and “Anchored offsets” in the pandas reference manual for details.

```
In [113]: fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6))
....: df1_hour.plot(ax=ax1, alpha=0.25)
....: df1_day.plot(ax=ax1)
....: df1_week.plot(ax=ax1)
....: df1_month.plot(ax=ax1)
....: df_diff.plot(ax=ax2)
....: ax2.set_title("temperature difference between outdoor and indoor")
....: fig.tight_layout()
```

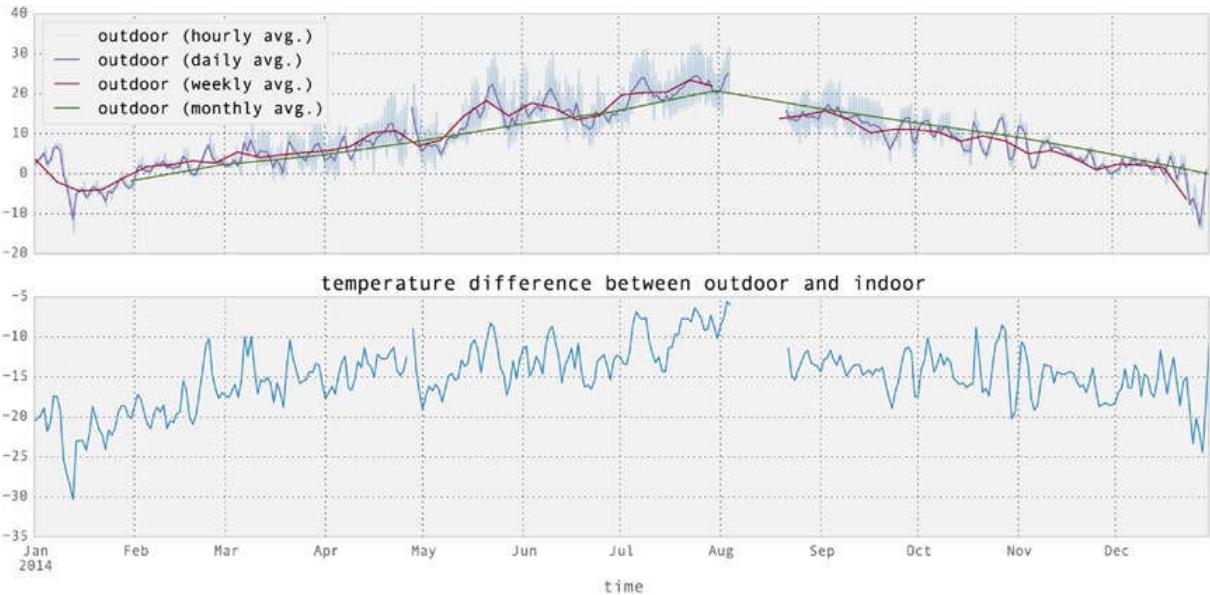


Figure 12-6. Outdoors temperature, resampled to hourly, daily, weekly, and monthly averages

As an illustration of upsampling, consider the following example where we resample the data frame `df1` to a sample frequency of 5 minutes, using three different fill methods (`None`, `ffill` for forward-fill, and `bfill` for back-fill). The original sample frequency is approximately 10 minutes, so this resampling is indeed upsampling. The result is three new data frames that we combine into a single `DataFrame` object using the `concat` function. The first five rows in the data frame are also shown below. Note that the every second data point is a new sample point, and depending on the value of the `fill_method` argument those values are filled (or not) according to the specified strategies. When no fill strategy is selected, the corresponding values are marked as missing using the `NaN` value.

```
In [114]: fill_methods = [None, 'ffill', 'bfill']
In [115]: pd.concat([df1.resample("5min", fill_method=fm).rename(columns={"outdoor": fm})
....:                 for fm in fill_methods], axis=1).head()
```

Out[115]:

	None	ffill	bfill
time			
2014-01-01 00:00:00+01:00	4.38	4.38	4.38
2014-01-01 00:05:00+01:00	NaN	4.38	4.25
2014-01-01 00:10:00+01:00	4.25	4.25	4.25
2014-01-01 00:15:00+01:00	NaN	4.25	4.19
2014-01-01 00:20:00+01:00	4.19	4.19	4.19

The Seaborn Graphics Library

The seaborn graphics library is built on top of Matplotlib, and it provides functions for generating graphs that are useful when working with statistics and data analysis, including distribution plots, kernel-density plots, joint distribution plots, factor plot, heat maps, facet plots, and several ways of visualizing regressions. It also provides methods for coloring data in graphs, and numerous well-crafted color palettes. The seaborn library is created with close attention to the aesthetics of the graphs it produces, and the graphs generated with the library tend to be both good looking and informative. The seaborn library distinguishes itself from the underlying Matplotlib library in that it provides polished higher-level graph functions for a specific application domain, namely, statistical analysis and data visualization. The ease with which standard statistical graphs can be generated with the library makes it a valuable tool in exploratory data analysis.

To get started using the seaborn library, we first import the `seaborn` module. Here we follow the common convention of importing this library under the name `sns`. After importing the library we can set a style for the graphs it produces using the `sns.set` function. Here we choose to work with the style called `darkgrid`, which produces graphs with a gray background (also try the `whitegrid` style).

```
In [116]: import seaborn as sns
In [117]: sns.set(style="darkgrid")
```

Importing `seaborn` and setting a style for the library alters the default settings for how Matplotlib graphs appear, including graphs produced by the `pandas` library. For example, consider the following plot of the previously used indoor and outdoor temperature time series. The resulting graph is shown in Figure 12-7, and although the graph was produced using the `pandas` `DataFrame` method `plot`, importing the `seaborn` library has changed the appearance of the graph (compare with Figure 12-3).

```
In [118]: df1 = pd.read_csv('temperature_outdoor_2014.tsv', delimiter='\t',
...:                         names=["time", "outdoor"])
...: df1.time = (pd.to_datetime(df1.time.values, unit="s")
...:             .tz_localize('UTC').tz_convert('Europe/Stockholm'))
...: df1 = df1.set_index("time").resample("10min")
In [119]: df2 = pd.read_csv('temperature_indoor_2014.tsv', delimiter='\t',
...:                         names=["time", "indoor"])
...: df2.time = (pd.to_datetime(df2.time.values, unit="s")
...:             .tz_localize('UTC').tz_convert('Europe/Stockholm'))
...: df2 = df2.set_index("time").resample("10min")
In [120]: df_temp = pd.concat([df1, df2], axis=1)
In [121]: fig, ax = plt.subplots(1, 1, figsize=(8, 4))
...: df_temp.resample("D").plot(y=["outdoor", "indoor"], ax=ax)
```

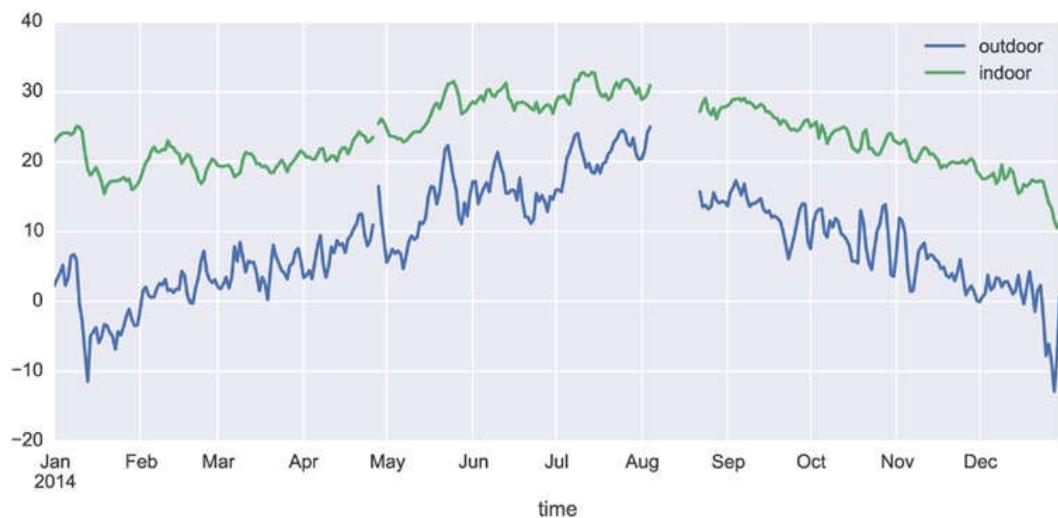


Figure 12-7. Time series plot produced by Matplotlib using the pandas library, with a plot style that is set up by the seaborn library

The main strength of the seaborn library, apart from generating good-looking graphics, is its collection of easy-to-use statistical plots. Examples of these are the `kdeplot` and `distplot`, which plot a kernel-density estimate plot and a histogram plot with a kernel-density estimate overlaid on top of the histogram, respectively. For example, the following two lines of code produce the graph shown in Figure 12-8. The solid blue and green lines in this figure are the kernel-density estimate that can also be graphed separately using the function `kdeplot` (not shown here).

```
In [122]: sns.distplot(df_temp.to_period("M")["outdoor"]["2014-04"].dropna().values, bins=50);
...: sns.distplot(df_temp.to_period("M")["indoor"]["2014-04"].dropna().values, bins=50);
```

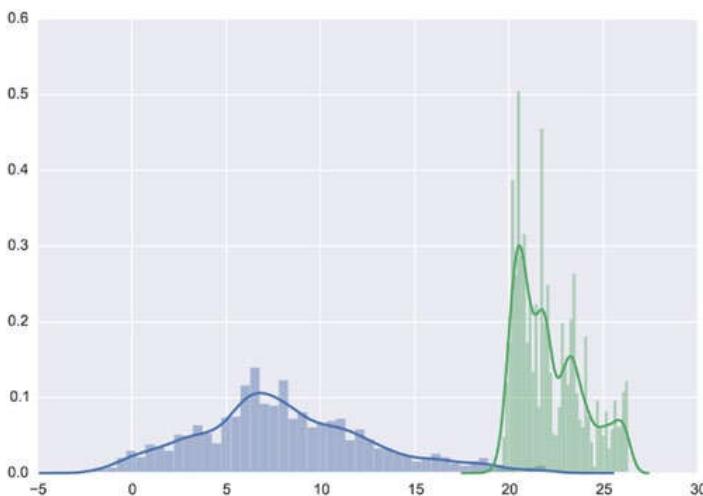


Figure 12-8. The histogram (bars) and kernel-density plots (solid lines) for the subset of the indoors and outdoors datasets that correspond to the month of April

The `kdeplot` function can also operate on two-dimensional data, showing a contour graph of the joint kernel-density estimate. Relatedly, we can use the `jointplot` function to plot the joint distribution for two separate datasets. Below we use the `kdeplot` and `jointplot` to show the correlation between the indoor and outdoor data series, which are resampled to hourly averages before visualized (we also drop missing values using `dropna` method, since the functions from the `seaborn` library do not accept arrays with missing data). The results are shown in Figure 12-9.

```
In [123]: sns.kdeplot(df_temp.resample("H")["outdoor"].dropna().values,
...:                 df_temp.resample("H")["indoor"].dropna().values, shade=False)
In [124]: with sns.axes_style("white"):
...:     sns.jointplot(df_temp.resample("H")["outdoor"].values,
...:                   df_temp.resample("H")["indoor"].values, kind="hex")
```

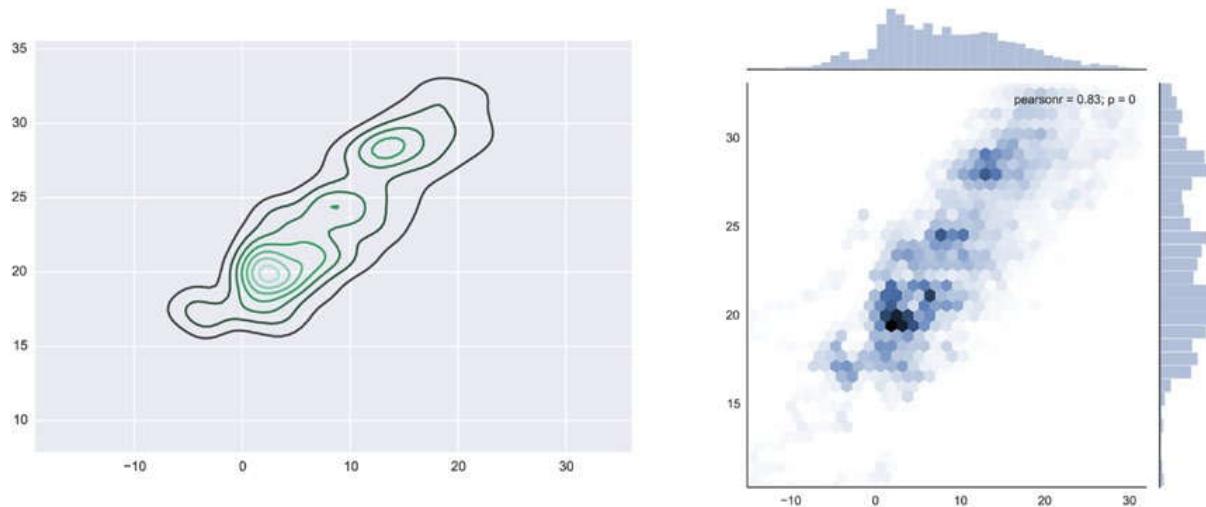


Figure 12-9. Two-dimensional kernel-density estimate contours (left) and the joint distribution for the indoor and outdoor temperature datasets (right). The outdoor temperatures are shown on the x-axis, and the indoor temperature on the y-axis

The `seaborn` library also provides functions for working with categorical data. A simple example of a graph type that is often useful for datasets with categorical variables is the standard boxplot for visualizing the descriptive statistics (min, max, median, and quartiles) of a dataset. An interesting twist on the standard boxplot is violin plot, in which the kernel-density estimated is shown in the width of boxplot. The `boxplot` and `violinplot` functions can be used to produce such graphs, as shown in the following example, and the resulting graph is shown in Figure 12-10.

```
In [125]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(8, 4))
...: sns.boxplot(df_temp.dropna(), ax=ax1, palette="pastel")
...: sns.violinplot(df_temp.dropna(), ax=ax2, palette="pastel")
```

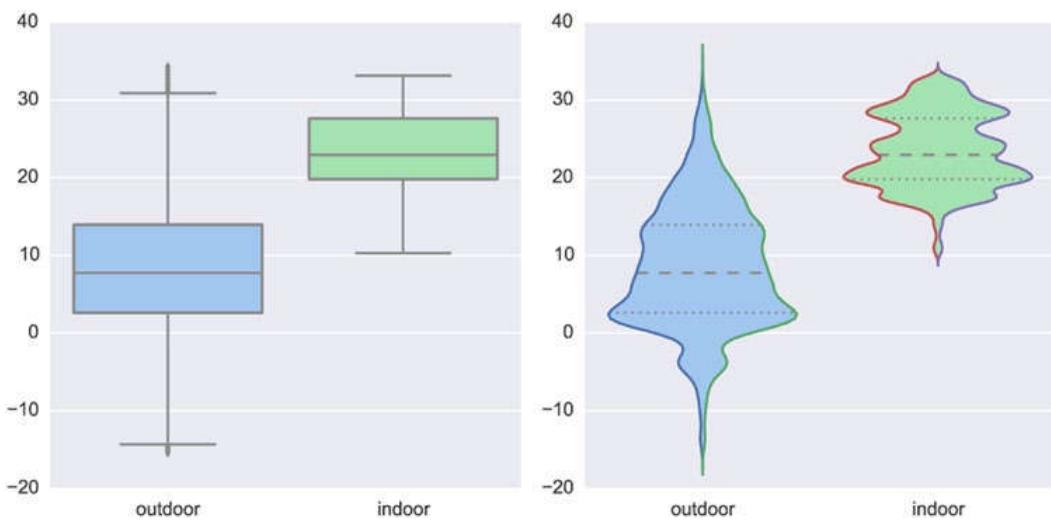


Figure 12-10. A box plot (left) and violin plot (right) for the indoor and outdoor temperature datasets

As a further example of violin plots, consider the outdoor temperature dataset partitioned by the month, which can be produced by passing the month field of the index of the data frame as second argument (used to group the data into categories). The resulting graph, which is shown in Figure 12-11, provides a compact and informative visualization of the distribution of temperatures for each month of the year.

```
In [126]: sns.violinplot(x=df_temp.dropna().index.month,
...:                      y=df_temp.dropna().outdoor, color="skyblue");
```

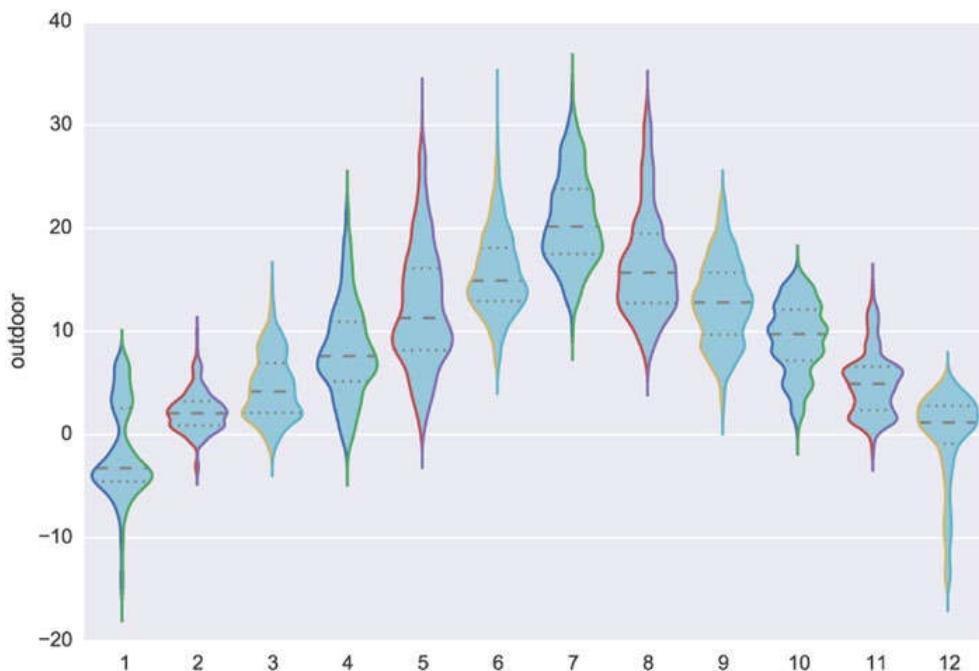


Figure 12-11. Violin plot for the outdoor temperature grouped by month

Heat maps are another type of graph that is handy when dealing with categorical variables, especially for variables with a large number of categories. The seaborn library provides the function `heatmap` for generating this type of graph. For example, working with the outdoors temperature data set, we can create two categorical columns `month` and `hour` by extracting those fields from the index and assign them to new columns in the data field. Next we can use the `pivot_table` function in pandas to pivot the columns into a table (matrix) where two selected categorical variables constitute the new index and columns. Here we pivot the temperature dataset so that the hours of the day are the columns, and the months of the year are the rows (index). To aggregate the multiple data points that falls within each hour-month category, we use `aggfunc=np.mean` argument to compute the mean of all the values:

```
In [127]: df_temp["month"] = df_temp.index.month
      ...: df_temp["hour"] = df_temp.index.hour
In [128]: table = pd.pivot_table(df_temp, values='outdoor', index=['month'],
      ...:                   columns=['hour'], aggfunc=np.mean)
```

Once we have created pivot table, we can visualize it as a heatmap using the `heatmap` function in seaborn. The result is shown in Figure 12-12.

```
In [129]: fig, ax = plt.subplots(1, 1, figsize=(8, 4))
      ...: sns.heatmap(table, ax=ax)
```

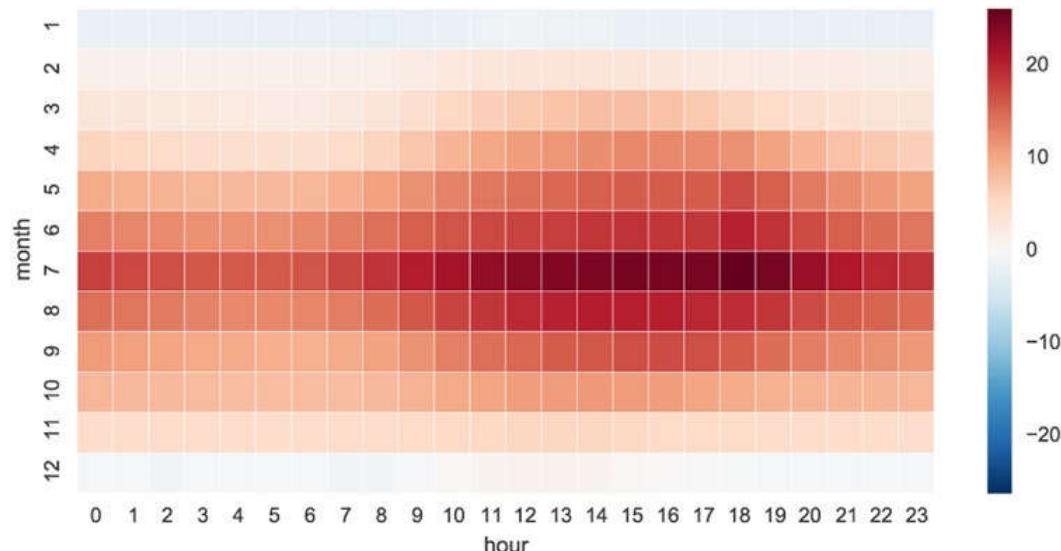


Figure 12-12. A heatmap of the outdoor temperature data grouped by hour of the day and month of the year

The seaborn library contains much more statistical visualization tools than what we have been able to survey here. However, I hope that looking at a few examples of what this library can do illustrates the essence of the seaborn library – that it is a convenient tool for statistical analysis and exploration of data, which is able to produce many standard statistical graphs with a minimal of effort. In the following chapters we will see further examples of applications of the seaborn library.

Summary

In this chapter we have explored data representation and data processing using the pandas library, and we briefly surveyed the statistical graphics tools provided by the seaborn visualization library. The pandas library provides the back end for much of data wrangling done with Python. It achieves this by adding a higher-level abstraction layer in the data representation on top of NumPy arrays, with additional methods for operating on the underlying data. The ease with which data can be loaded, transformed, and manipulated makes it an invaluable part of the data processing workflow in Python. The pandas library also contains basic functions for visualizing the data that is represented by its data structures. Being able to quickly visualize data represented as pandas series and data frames is an important tool in exploratory data analytics as well as for presentation. The seaborn library takes this a step further, and provides a rich collection of statistical graphs that can be produced often with a single line of code. Many functions in the seaborn library can operate directly on pandas data structures.

Further Reading

A great introduction to the pandas library is given by the original creator of the library in (McKinney, 2013), and it is also a rather detailed introduction to NumPy. The pandas official documentation, available at <http://pandas.pydata.org/pandas-docs/stable>, also provides an accessible and very detailed description of the features of the library. Another good online resource for learning pandas is <http://github.com/jvns/pandas-cookbook>. For data visualization, we have looked at the seaborn library in this chapter, and it is well described in the documentation available on its web site. With respect to higher-level visualization tools, it is also worth exploring the ggplot library for Python: <http://ggplot.yhatq.com>, which is an implementation based on renowned *Grammar of graphics* (L. Wilkinson, 2005). This library is also closely integrated with the pandas library, and it provides statistical visualization tools that are convenient when analyzing data. For more information about visualization in Python, see for example, the book by Vaingast.

References

- McKinney, W. (2013). *Python for Data Analysis*. Sebastopol: O'Reilly.
Vaingast, S. (2014). *Beginning Python Visualization*. New York: Apress.
L. Wilkinson, D. W. (2005). *The Grammar of Graphics*. Chicago: Springer.